



SINGAPORE
INSTITUTE OF
TECHNOLOGY

Introduction to Recommender Systems and Collaborative Filtering

Lucas Vinh Tran

16 March 2023

About Me

- Principal Scientist / Director of Personalization at JPMorgan Chase; Previous at Apple
- Lead Personalization and Recommendation team
- Before:
 - Ph.D. in Computer Science (Personalized and Group Recommender Systems)

Our Aims

- It is not possible to cover every aspect of the recommender systems in depth within 3 lectures
- We will try to cover a broad spectrum of recommender systems, from traditional algorithms such as collaborative filtering, matrix factorization, deep learning, ... to recommender systems in industry

Lectures

- Week 11: Introduction to Recommender Systems and Collaborative Filtering
- Week 12: Recommender Systems before Deep Learning era
 - Lab: Focus on topics in Week 11
- Week 13: Modern Recommender Systems in Theory and Practice
 - Lab: Focus on topics in Week 12 and 13

Introduction

What is a Recommender System?

Example - Netflix

The first screenshot shows the main Netflix homepage with sections for "Blockbuster Sci-Fi & Fantasy", "Trending now", and "Continue watching for". The second screenshot shows a movie detail page for "Slick" with a "97% match" rating.

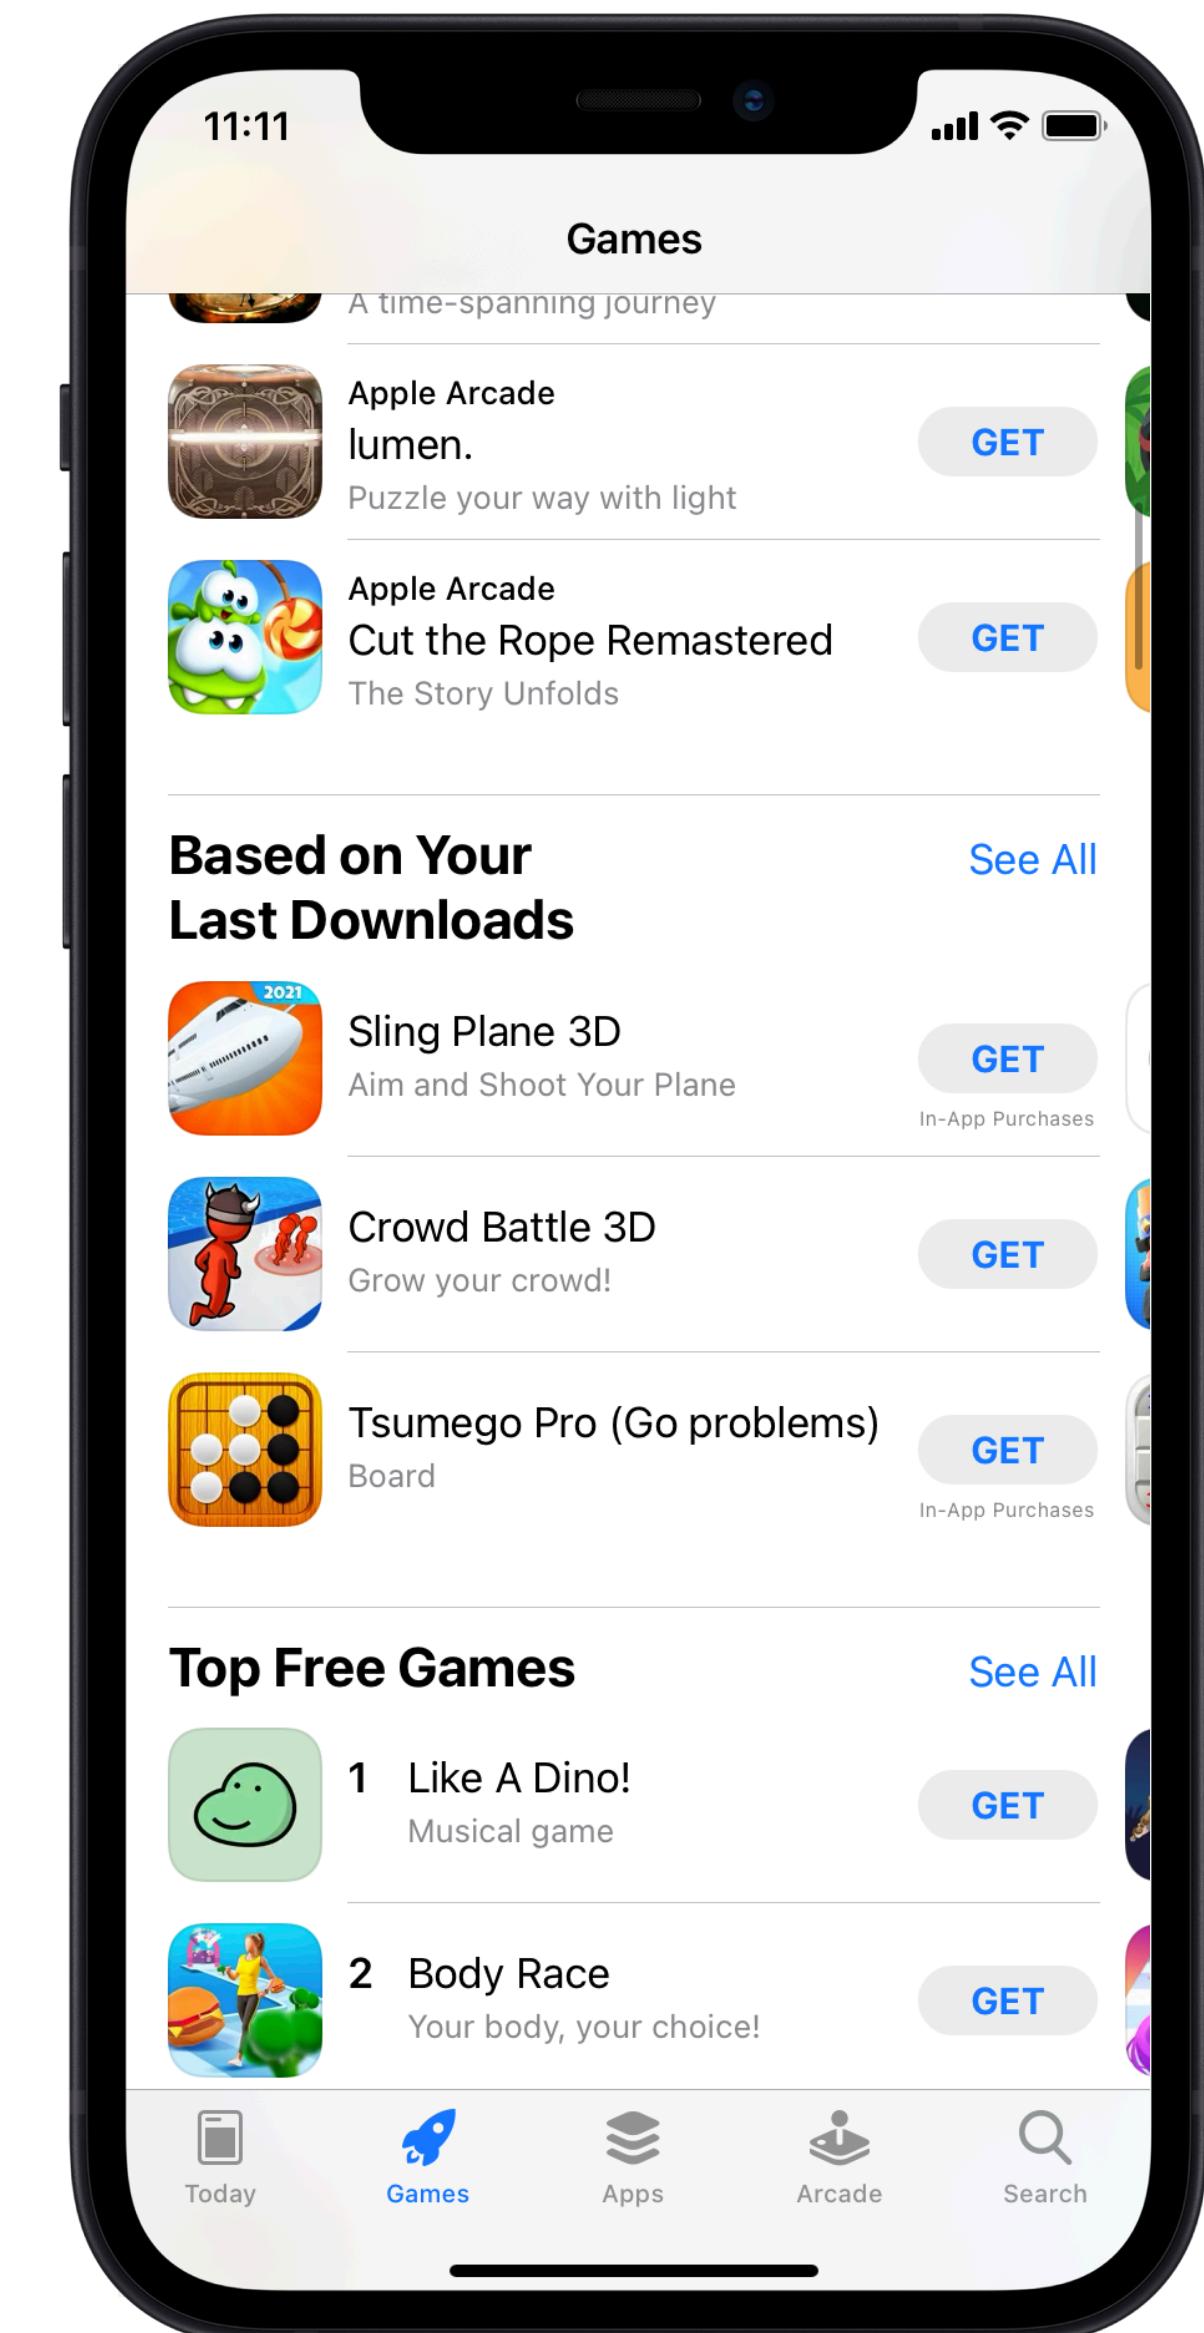
Frequently Bought Together

Price for all three: \$74.20

- This item: Beginning Ruby: From Novice to Professional (Expert's Voice in Open Source) by Peter Cooper Paperback \$27.78
- Learn to Program, Second Edition (The Facets of Ruby Series) by Chris Pine Paperback \$16.94
- Ruby on Rails Tutorial: Learn Web Development with Rails (2nd Edition) (Addison-Wesley Professional Ruby ... by Michael Hartl Paperback \$29.48

Customers Who Bought This Item Also Bought

 Learn to Program, Second Edition (The Facets of... Chris Pine ★★★★★ 42 Paperback \$16.94 ✓Prime	 The Well-Grounded Rubyist David A. Black ★★★★★ 39 Paperback \$32.49 ✓Prime	 Ruby on Rails Tutorial: Learn Web Development... Michael Hartl ★★★★★ 70 Paperback \$29.48 ✓Prime	 The Ruby Programming Language David Flanagan ★★★★★ 74 Paperback \$26.35 ✓Prime	 The Well-Grounded Rubyist David A. Black ★★★★★ 19 #1 Best Seller in Ruby Programming Computer Paperback \$29.67 ✓Prime
--	---	---	---	--

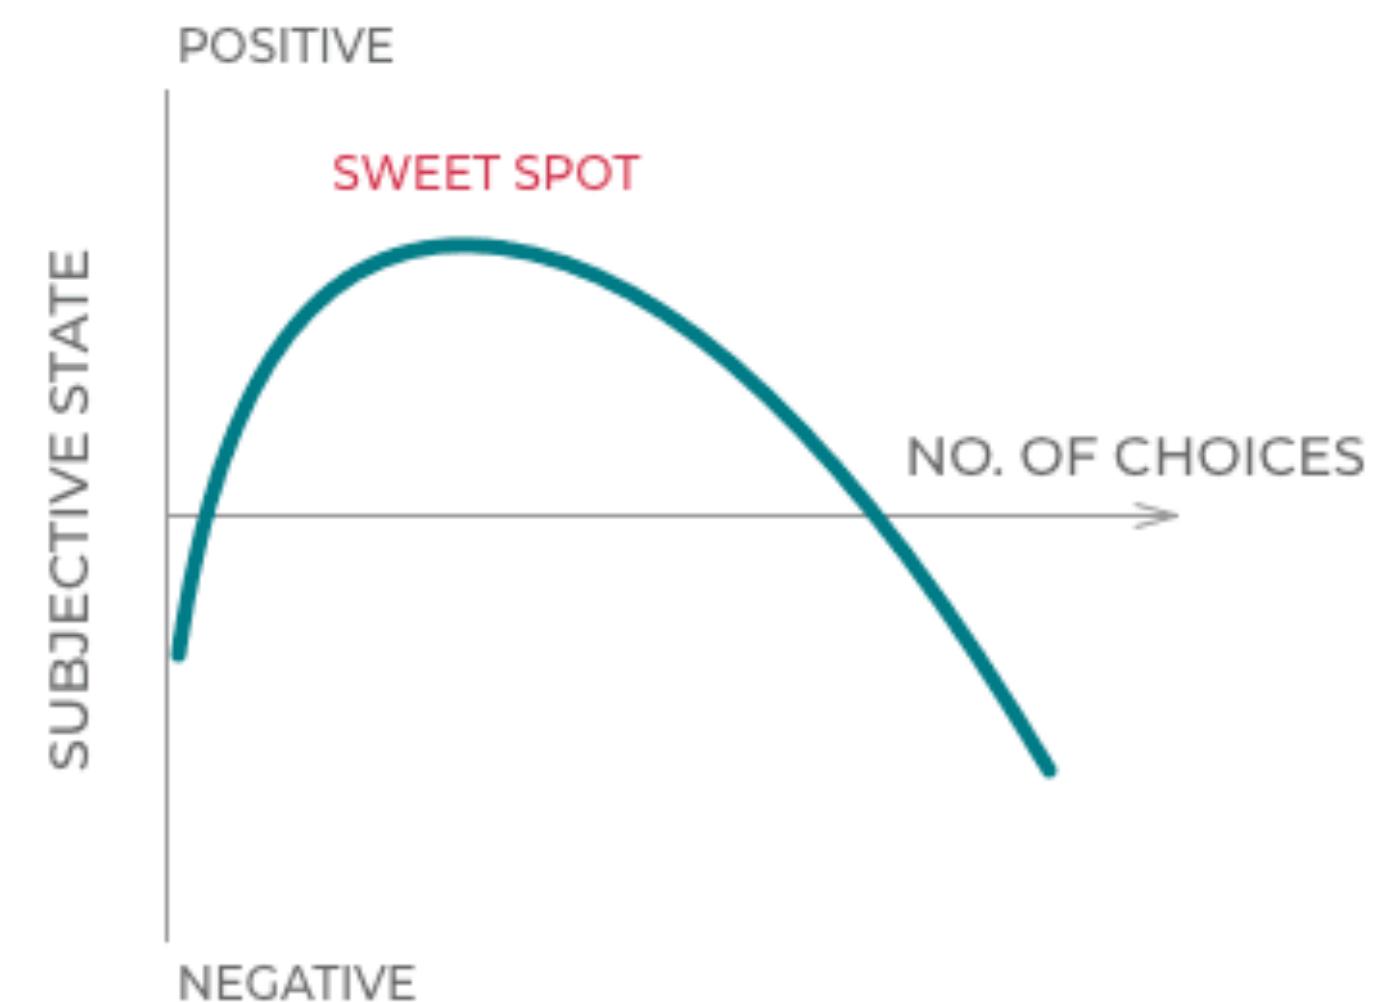


Why recommender systems?

- Netflix: 2/3 of the watched movies are recommended
- Google News: recommendations generate 38% more clickthrough
- Amazon: 35% sales achieved are from recommendations
- ...



"Can I tell you about a few items that aren't on the menu?"



Why recommender systems?

Bring values to **customers**

- Find things that are more relevant
- Narrow down the set of choices
- Explore and discover more interesting spaces
- ...

Why recommender systems?

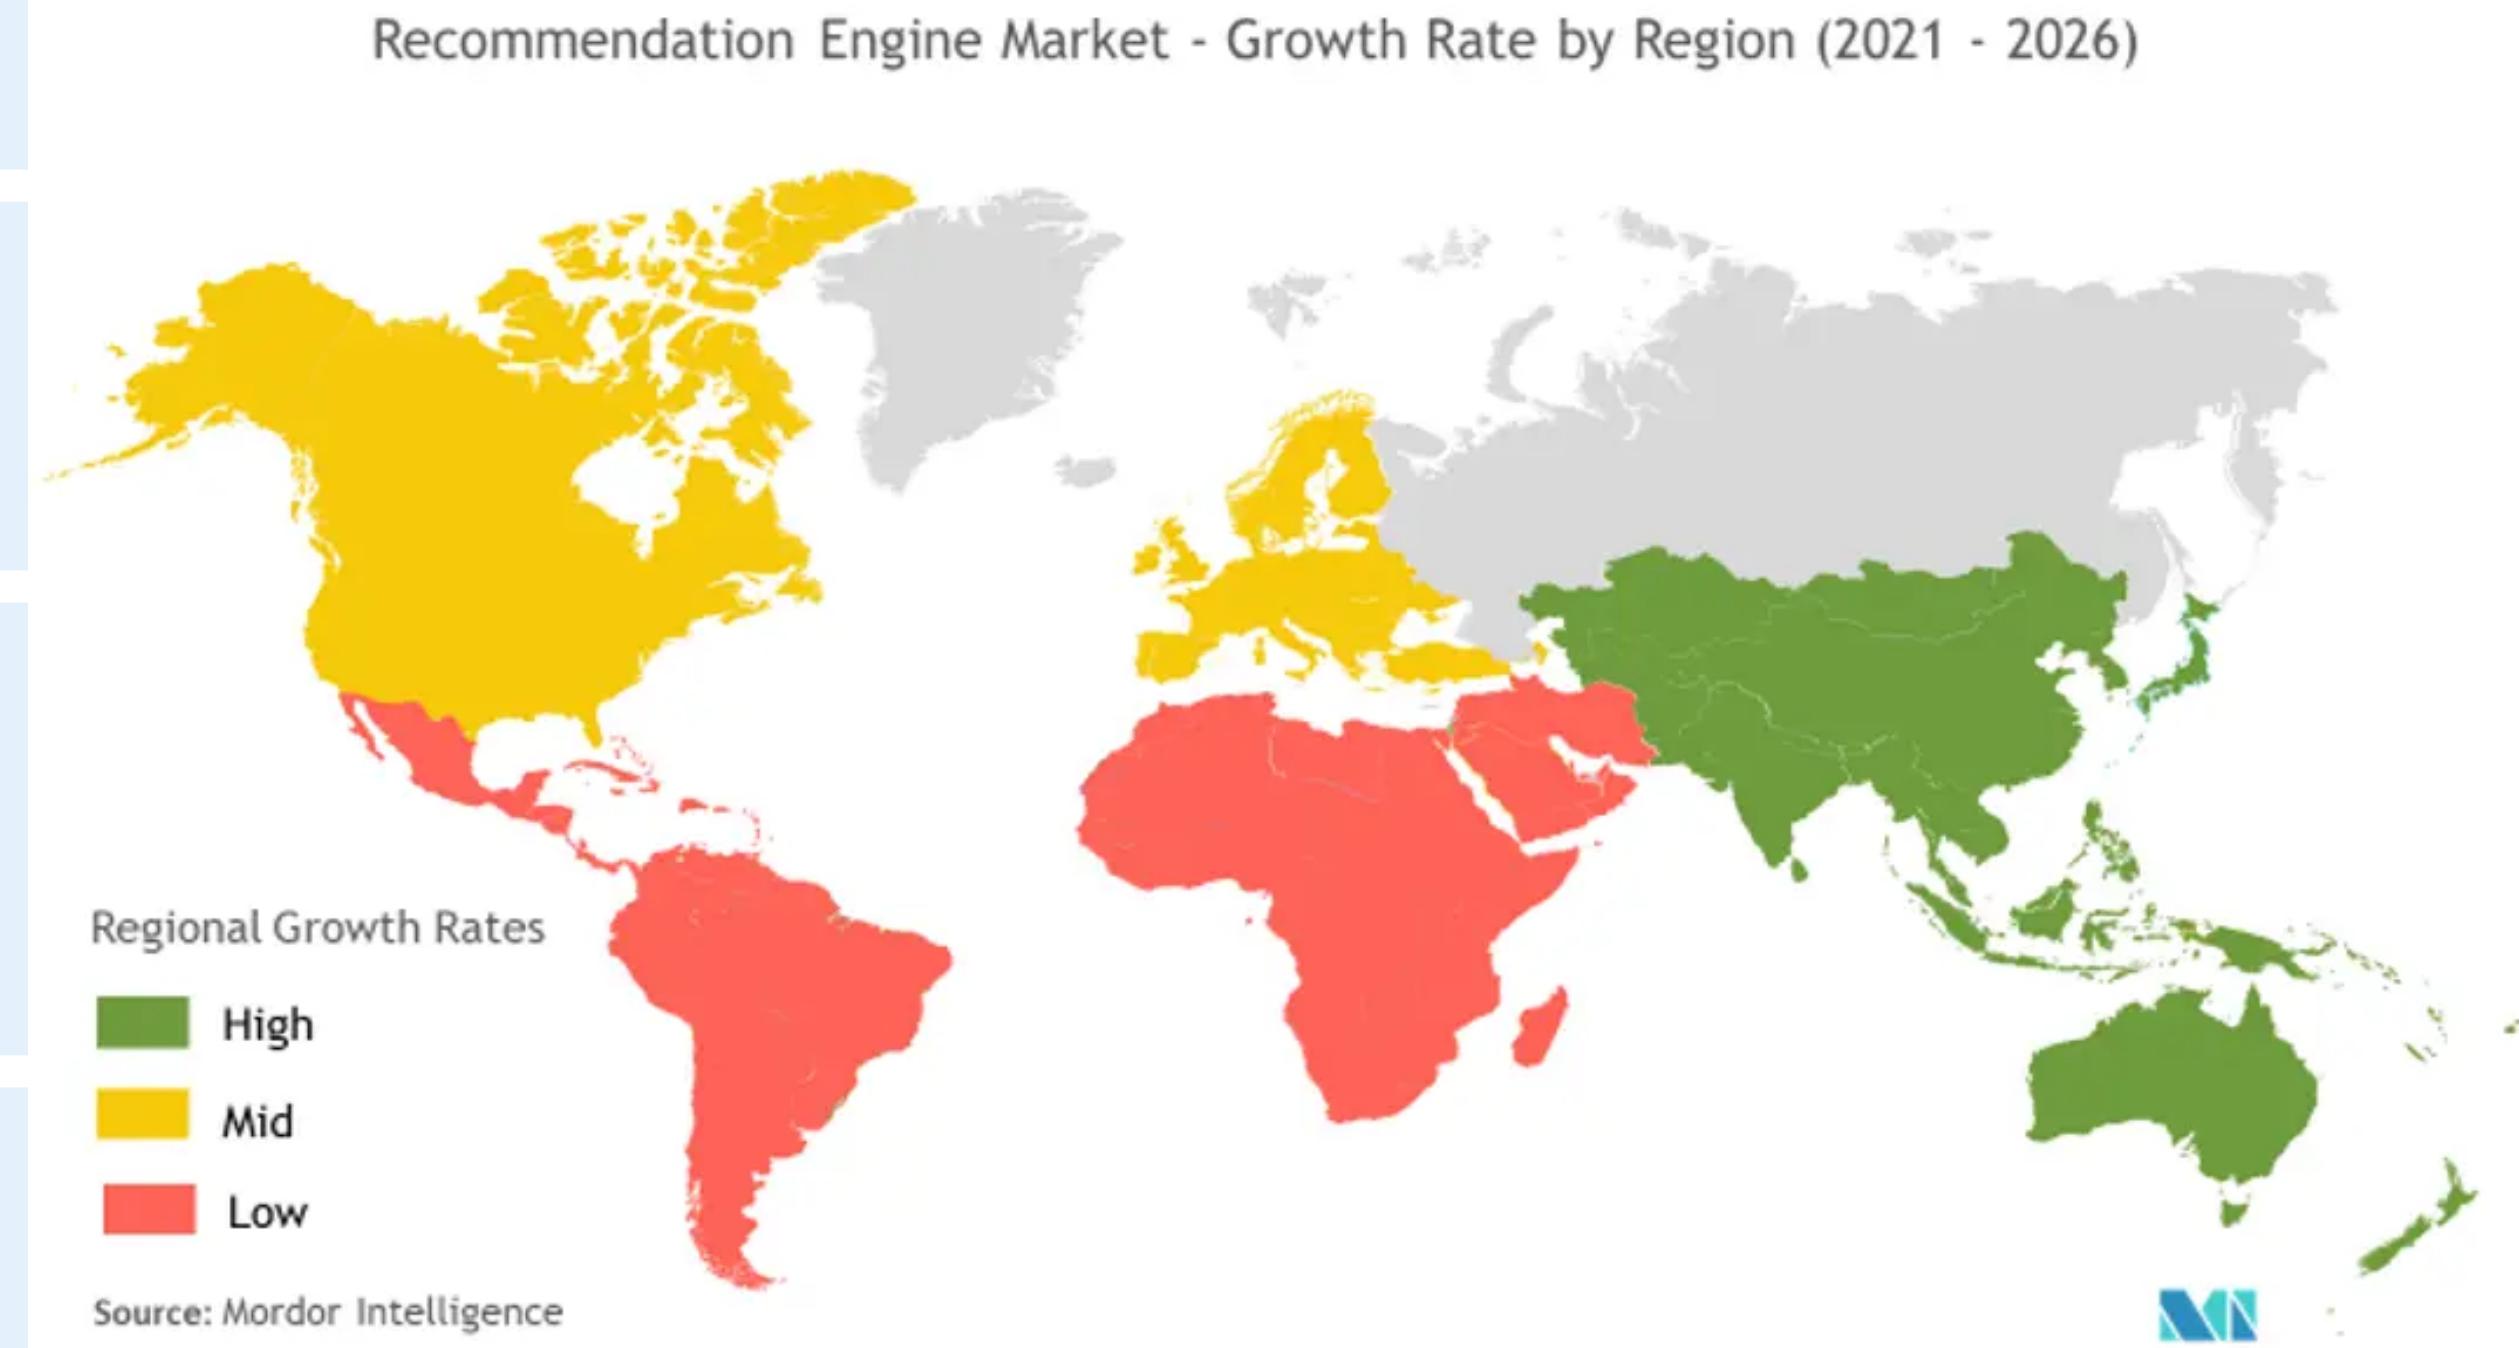
Bring values to **providers**

- Improve profit by providing relevant and accurate recommendations
- Increase customer retention (e.g., customer loyalty)
- Guide consuming behaviours
- ...

Everything is Recommendation!

- The Recommendation Engine market was valued at USD 2.12 billion in 2020, and it is expected to reach USD 15.13 billion by 2026

Deployment Mode	On-Premise Cloud
Types	Collaborative Filtering Content-Based Filtering Hybrid Recommendation Systems Other Types
End-user Industry	IT and Telecommunication BFSI Retail Media and Entertainment Healthcare Other End-user Industries
Geography	North America Europe Asia Pacific Latin America Middle East and Africa



Example

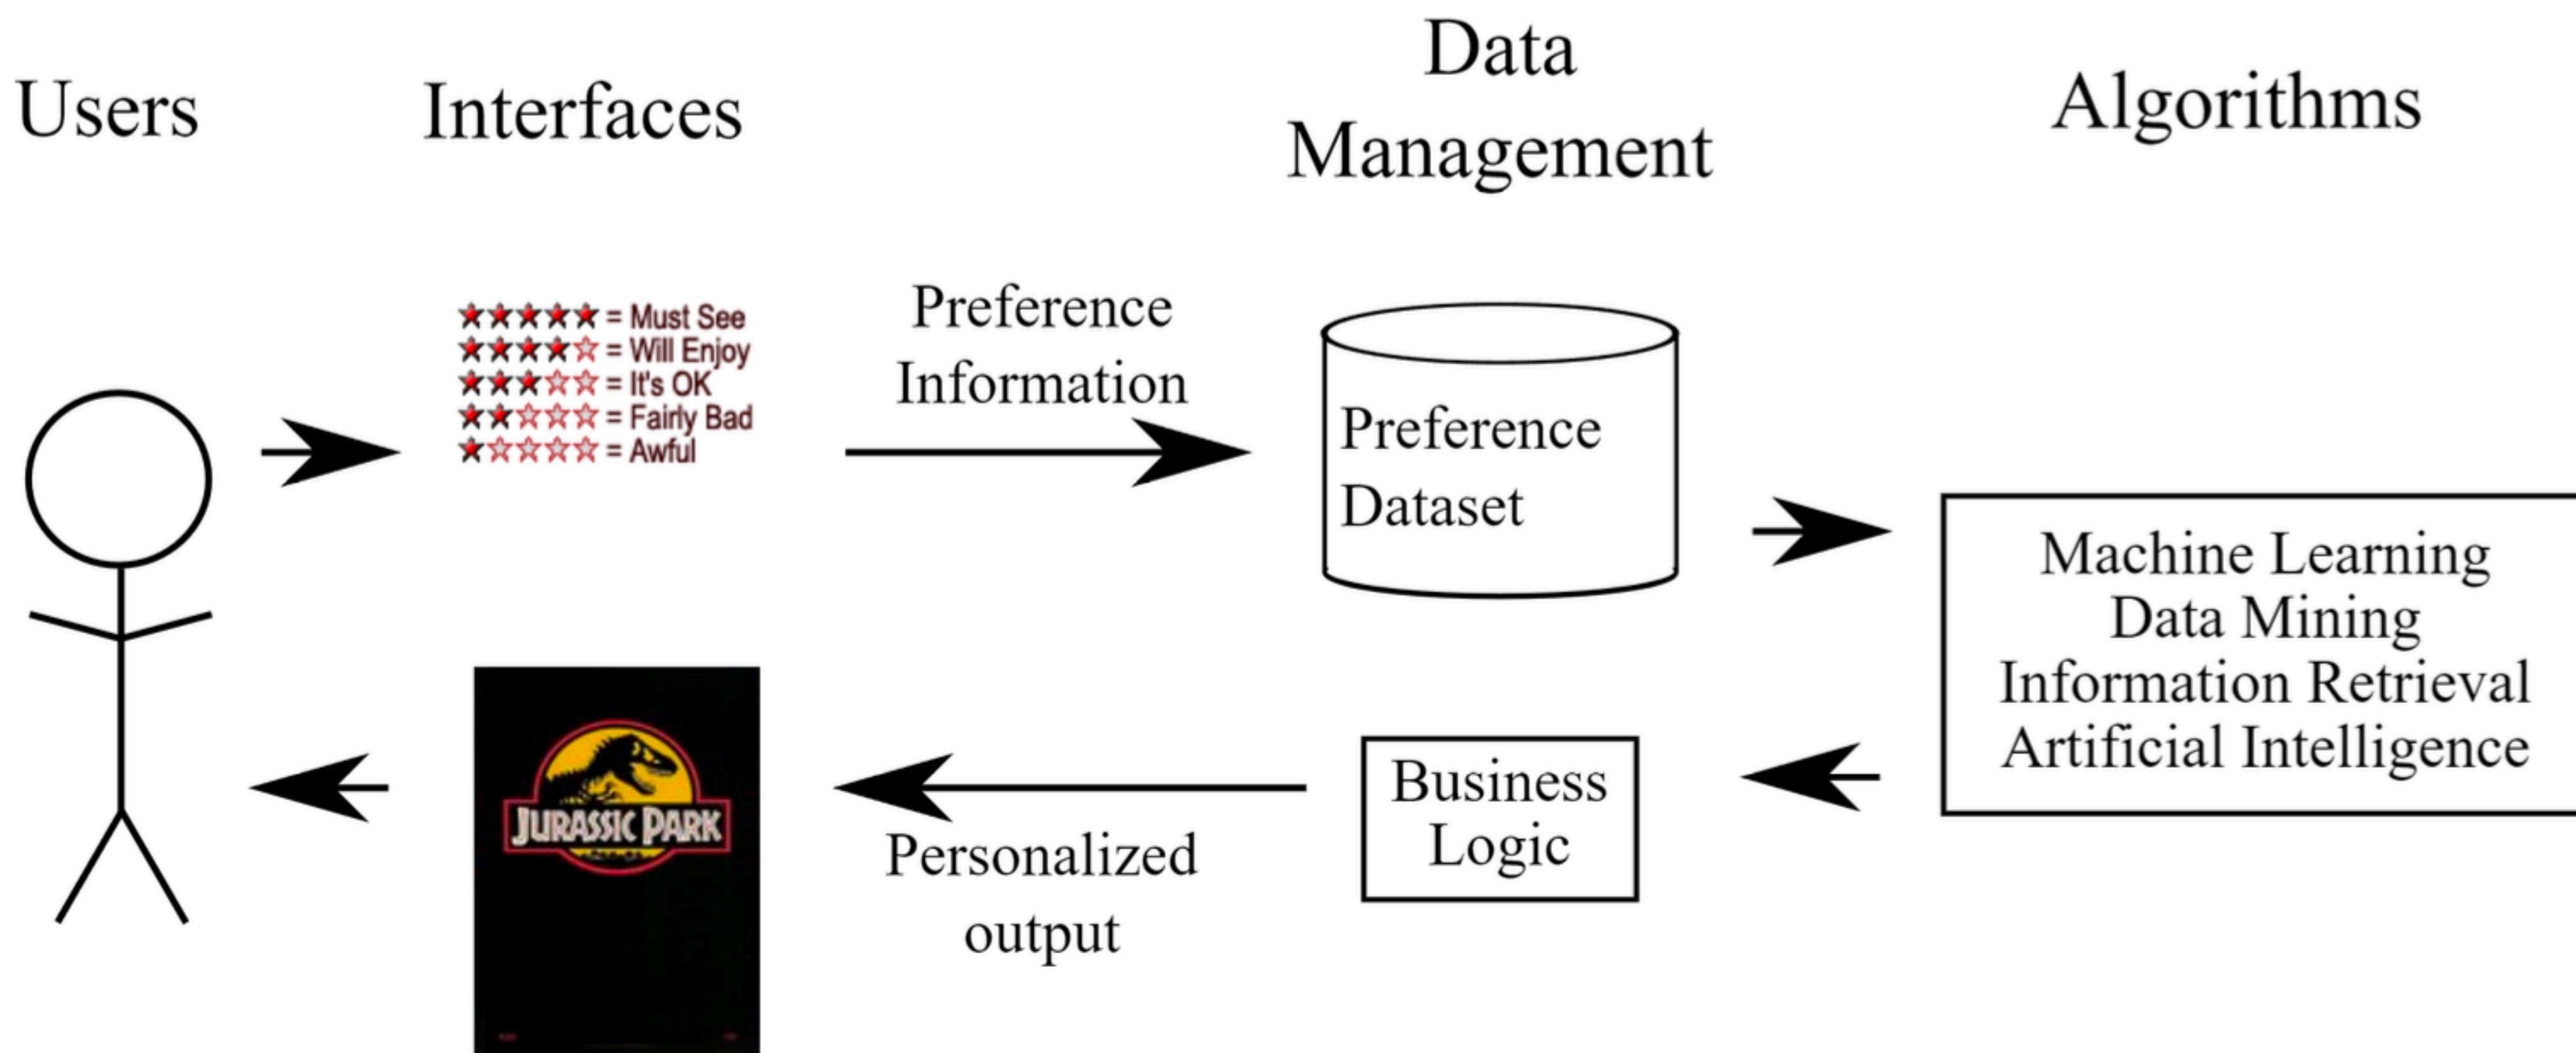
According to TikTok: “The system recommends content by ranking videos based on a combination of factors – starting from interests you express as a new user and adjusting for things you indicate you’re not interested in, too.”

- User Activity: likes, shares, comments, ...
- Content: video, audio, hashtags, effects, ...
- Auxiliary information: gender, country, language, ...
- Social information: follower/followee, tags, linked to other accounts, ...
- Contextual information: location, geographic, time, ...



Pipeline of RS

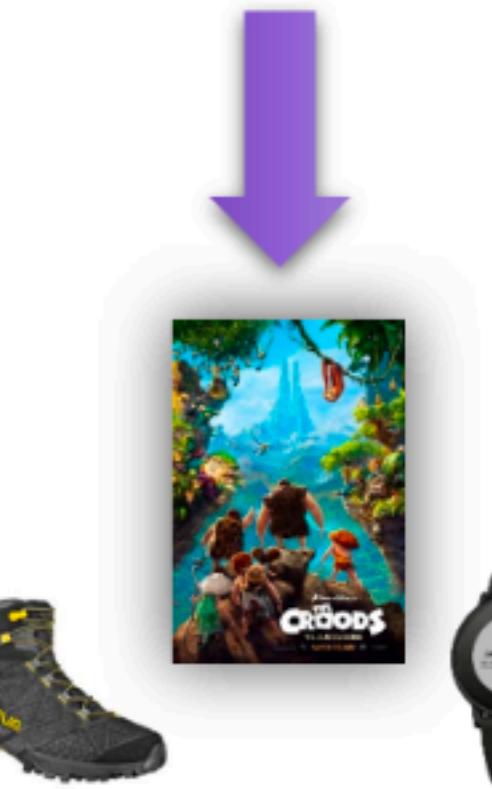
- Learn a utility function that predicts a user's preference towards an item



Data Source of RS

- Users: structured, semi-structured, unstructured
- Items: structured, semi-structured, unstructured
- Interactions: ratings, like/dislike
- Contexts: side information
- Knowledge: domain experts
- ...

Formulation of RS



“RS are software agents that elicit the interests and preferences of individual consumers and make recommendations accordingly.”

“They have the potential to support and improve the quality of the decisions consumers make while searching for and selecting products online.”

Goal:

Learn a ***utility function*** that ***predicts*** a user's preference towards an item

❑ Inputs:

- User model (e.g. implicit/explicit feedbacks, preferences, demographics, social connections, situational context)
- Items (with or without description of item characteristics)

❑ Outputs:

- Predicted preference scores (Used for ranking)

Challenges of RS

- Scalability
- Sparsity/Density
- Cold-start
- Biases and Fairness
- Privacy and Security
- ...

Collaborative Filtering

Collaborative Filtering

Basic Idea

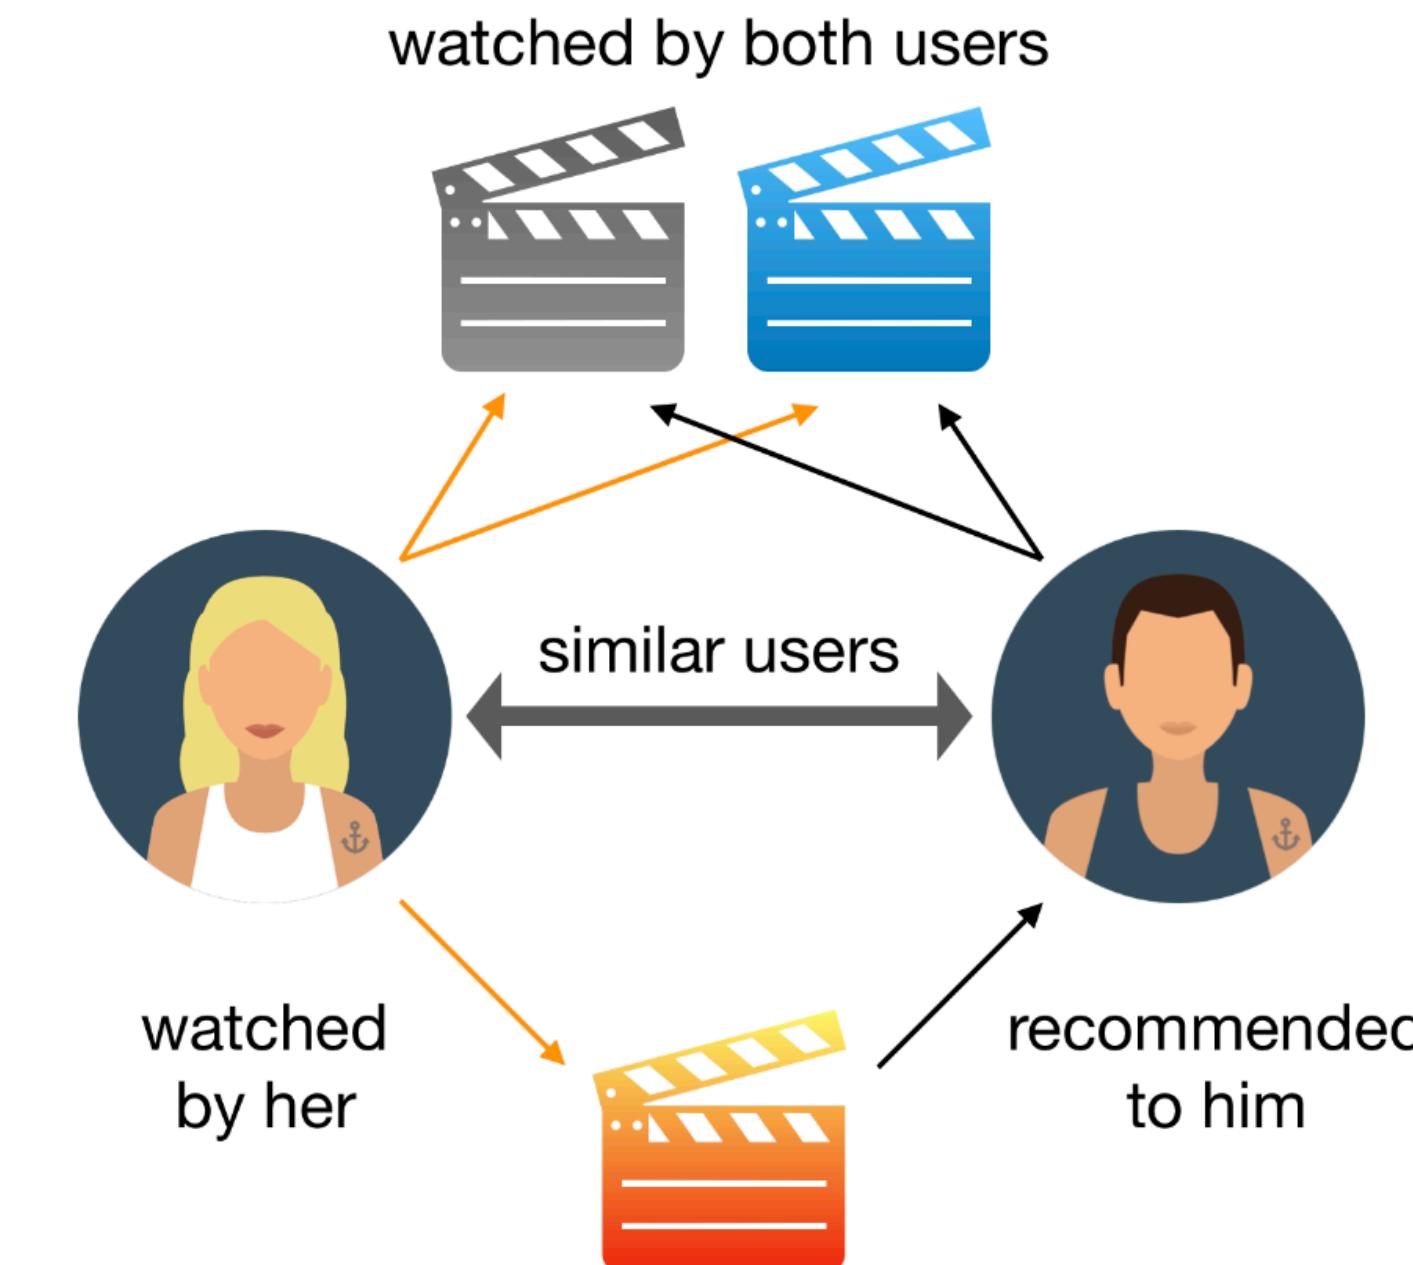
- “Wisdom of the crowd” – look at the ratings of like-minded users to recommend items
- Analyze the *known* preferences of a group of users to make predictions of the unknown preferences for other users.



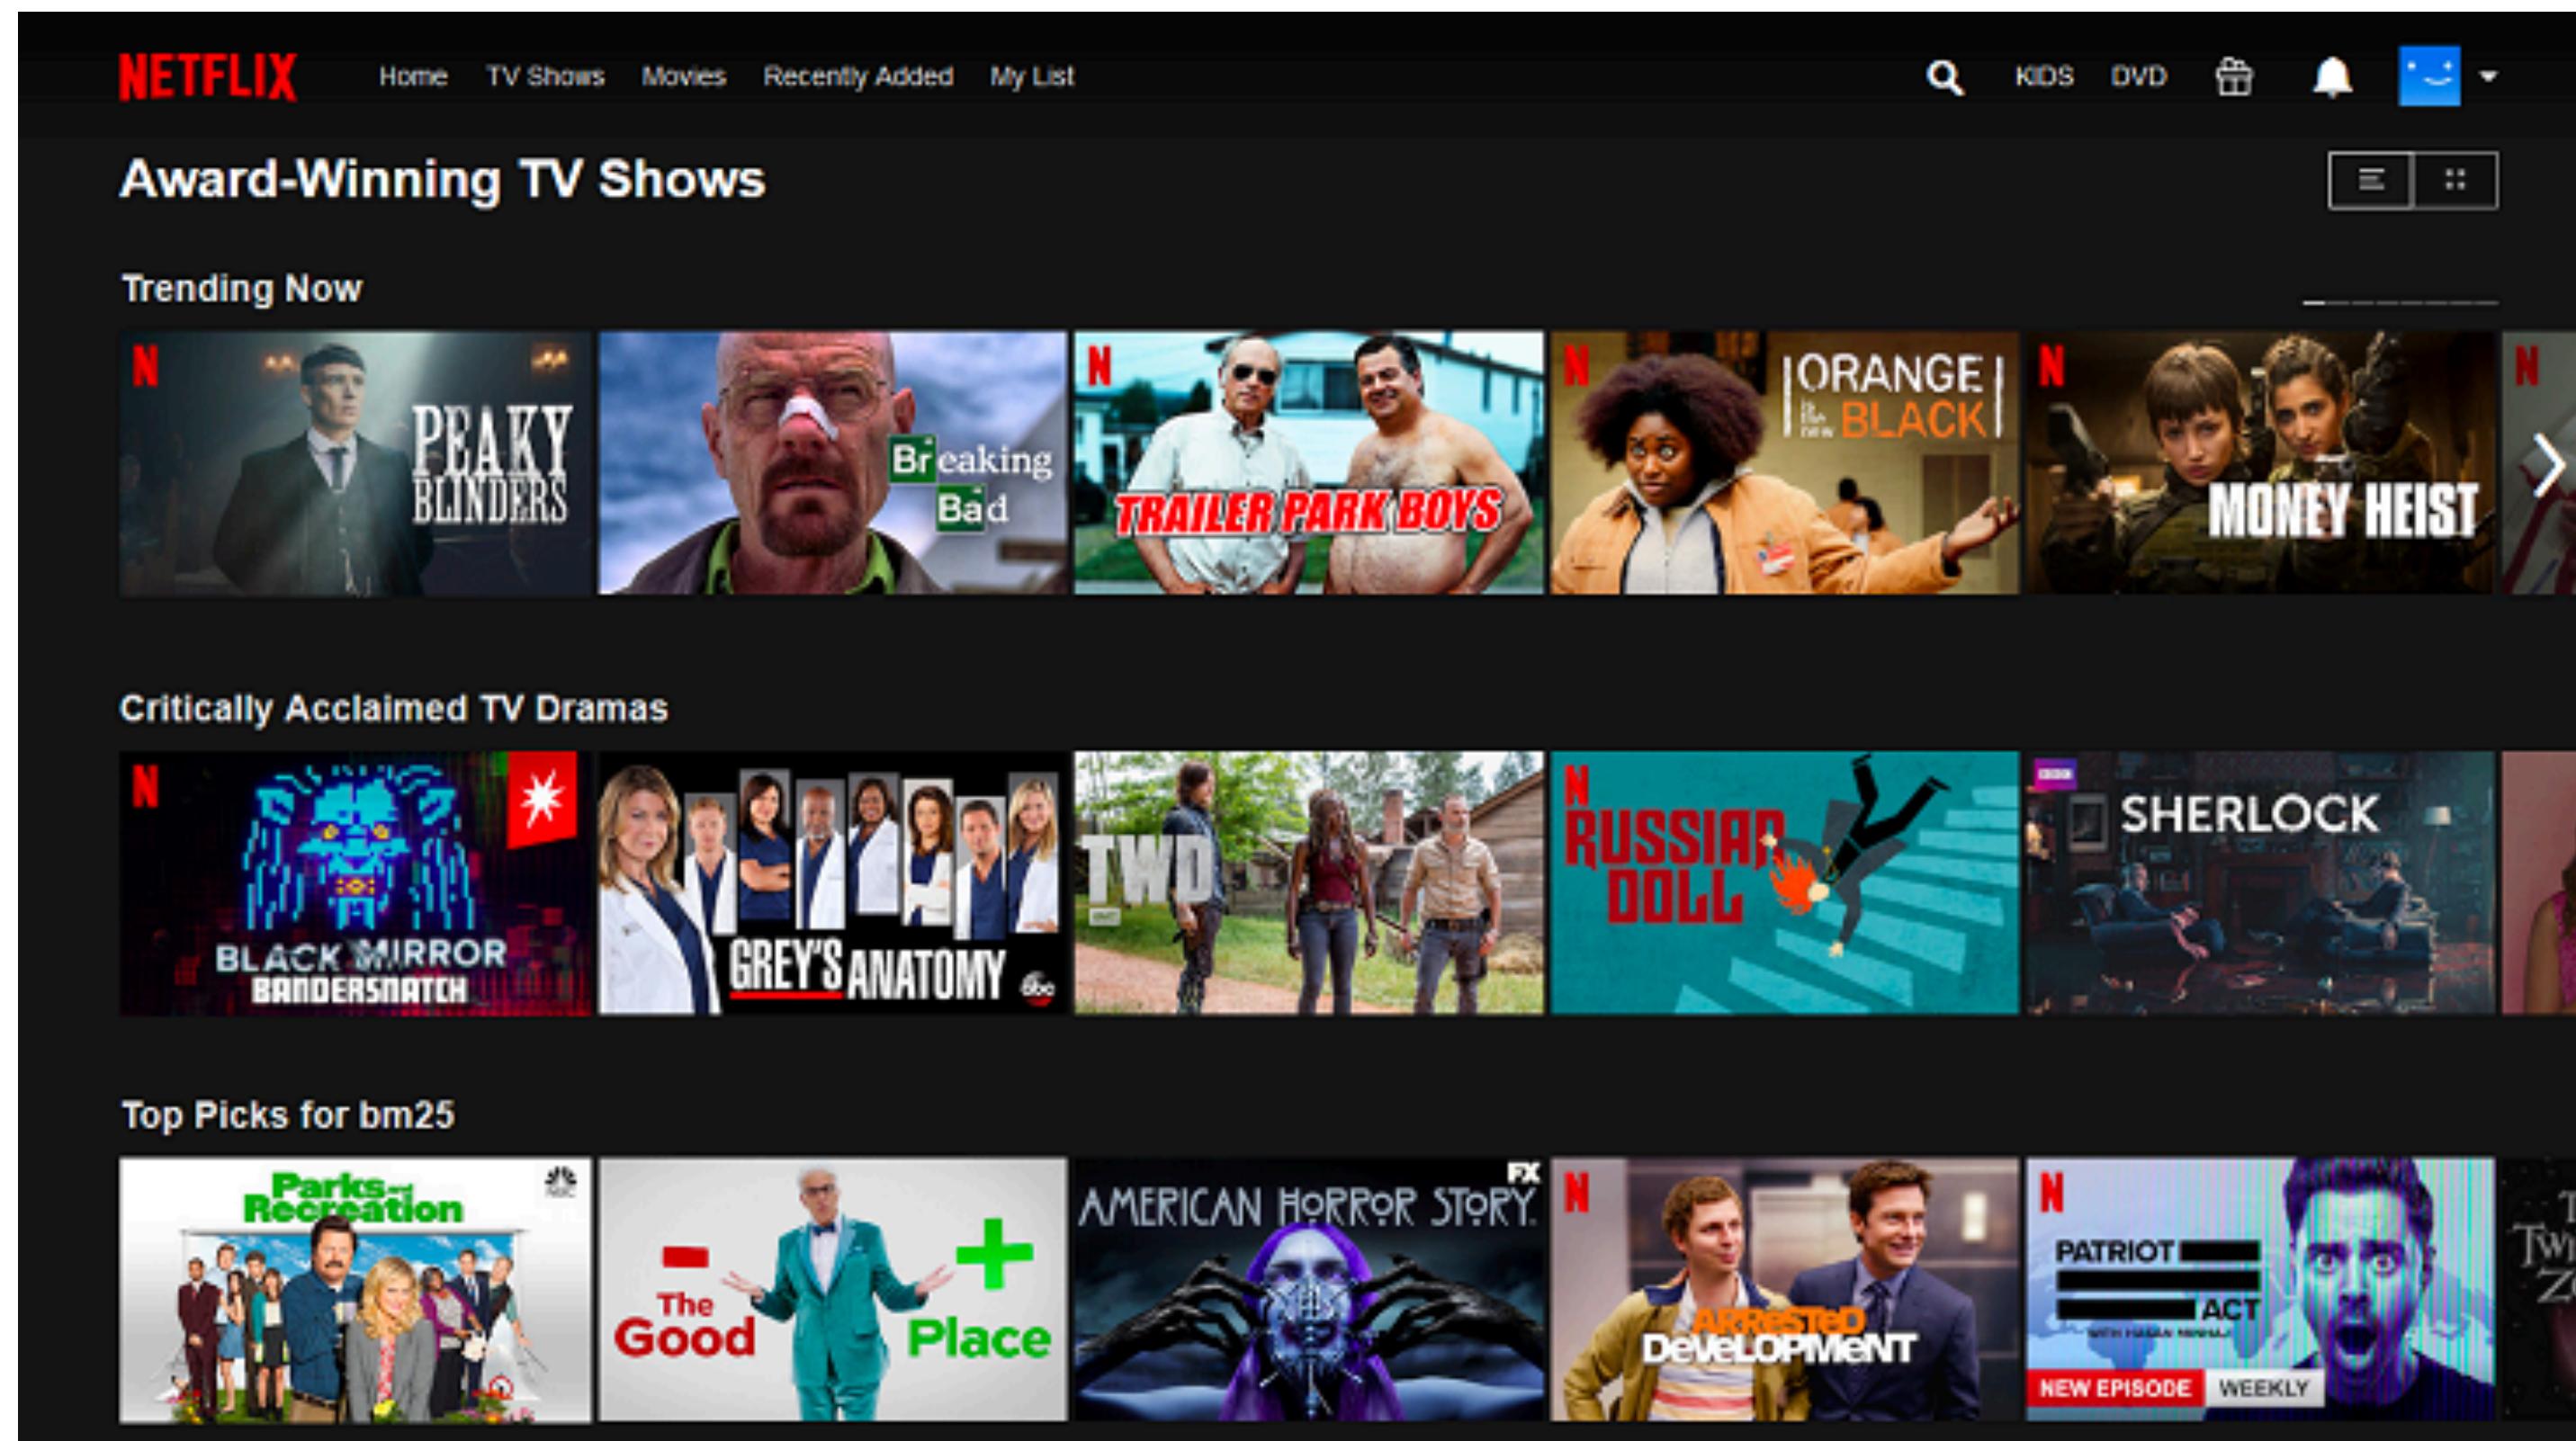
Collaborative Filtering

Basic Assumptions

- Users give ratings to a group of items
- Users who had similar tastes in the past, will have similar tastes in the future



Collaborative Filtering



	Breaking Bad	American Horror Story	Money Heist	Sherlock	Grey's Anatomy
Lucas	5	3	4	4	?
User 1	3	1	2	3	3
User 2	4	3	4	3	5
User 3	3	3	1	5	4
User 4	1	5	5	2	1

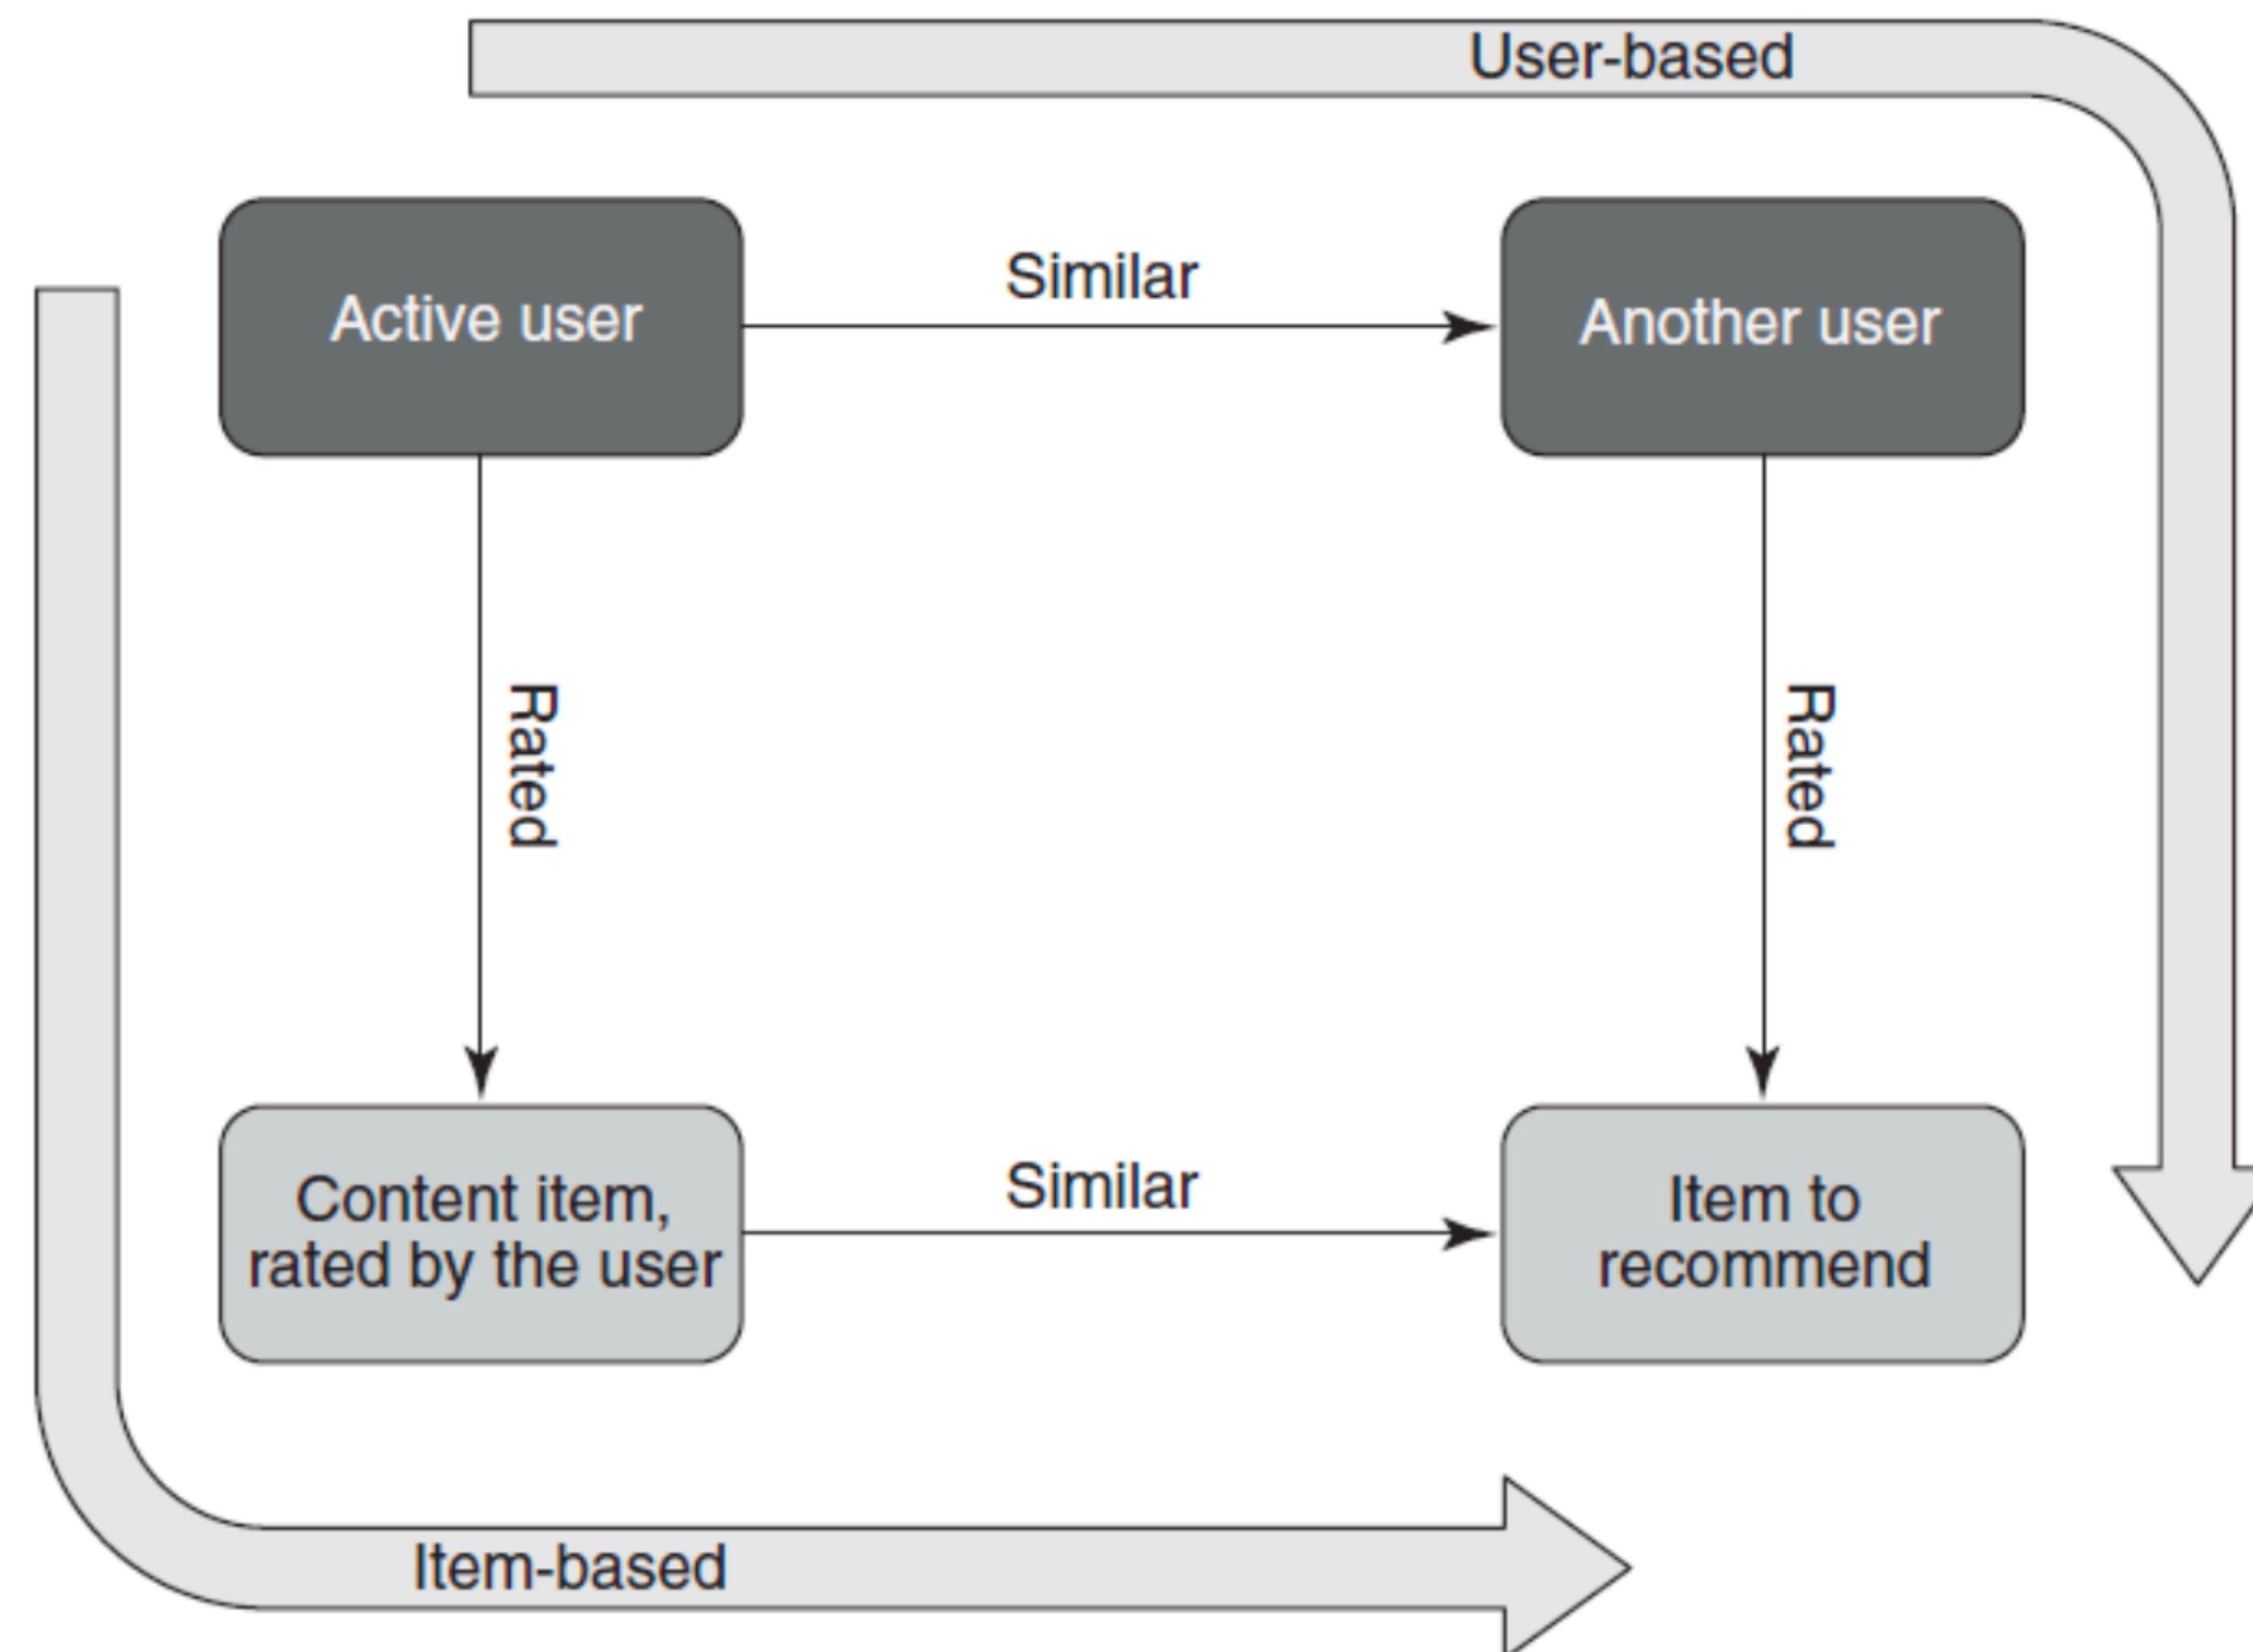
Collaborative Filtering

- Neighbourhood-based: Compute similarities between users or items (i.e., “memory” of the system) that are successively exploited to produce recommendations
- User-based: finding similar users
- Item-based: finding similar items
- Model-based: Estimate or learn a model, then apply this model to make rating predictions

Neighbourhood-based CF

Neighbourhood-based CF

User-based vs. Item-based



User-based CF

- If users had similar tastes in the past, they will have similar tastes in the future
- User preferences retain stable and consistent over time

	Grey's Anatomy	American Horror Story	Money Heist	Sherlock	Breaking Bad
Lucas	5	3	4	4	?
User 1	3	1	2	3	3
User 2	4	3	4	3	5
User 3	3	3	1	5	4
User 4	1	5	5	2	1

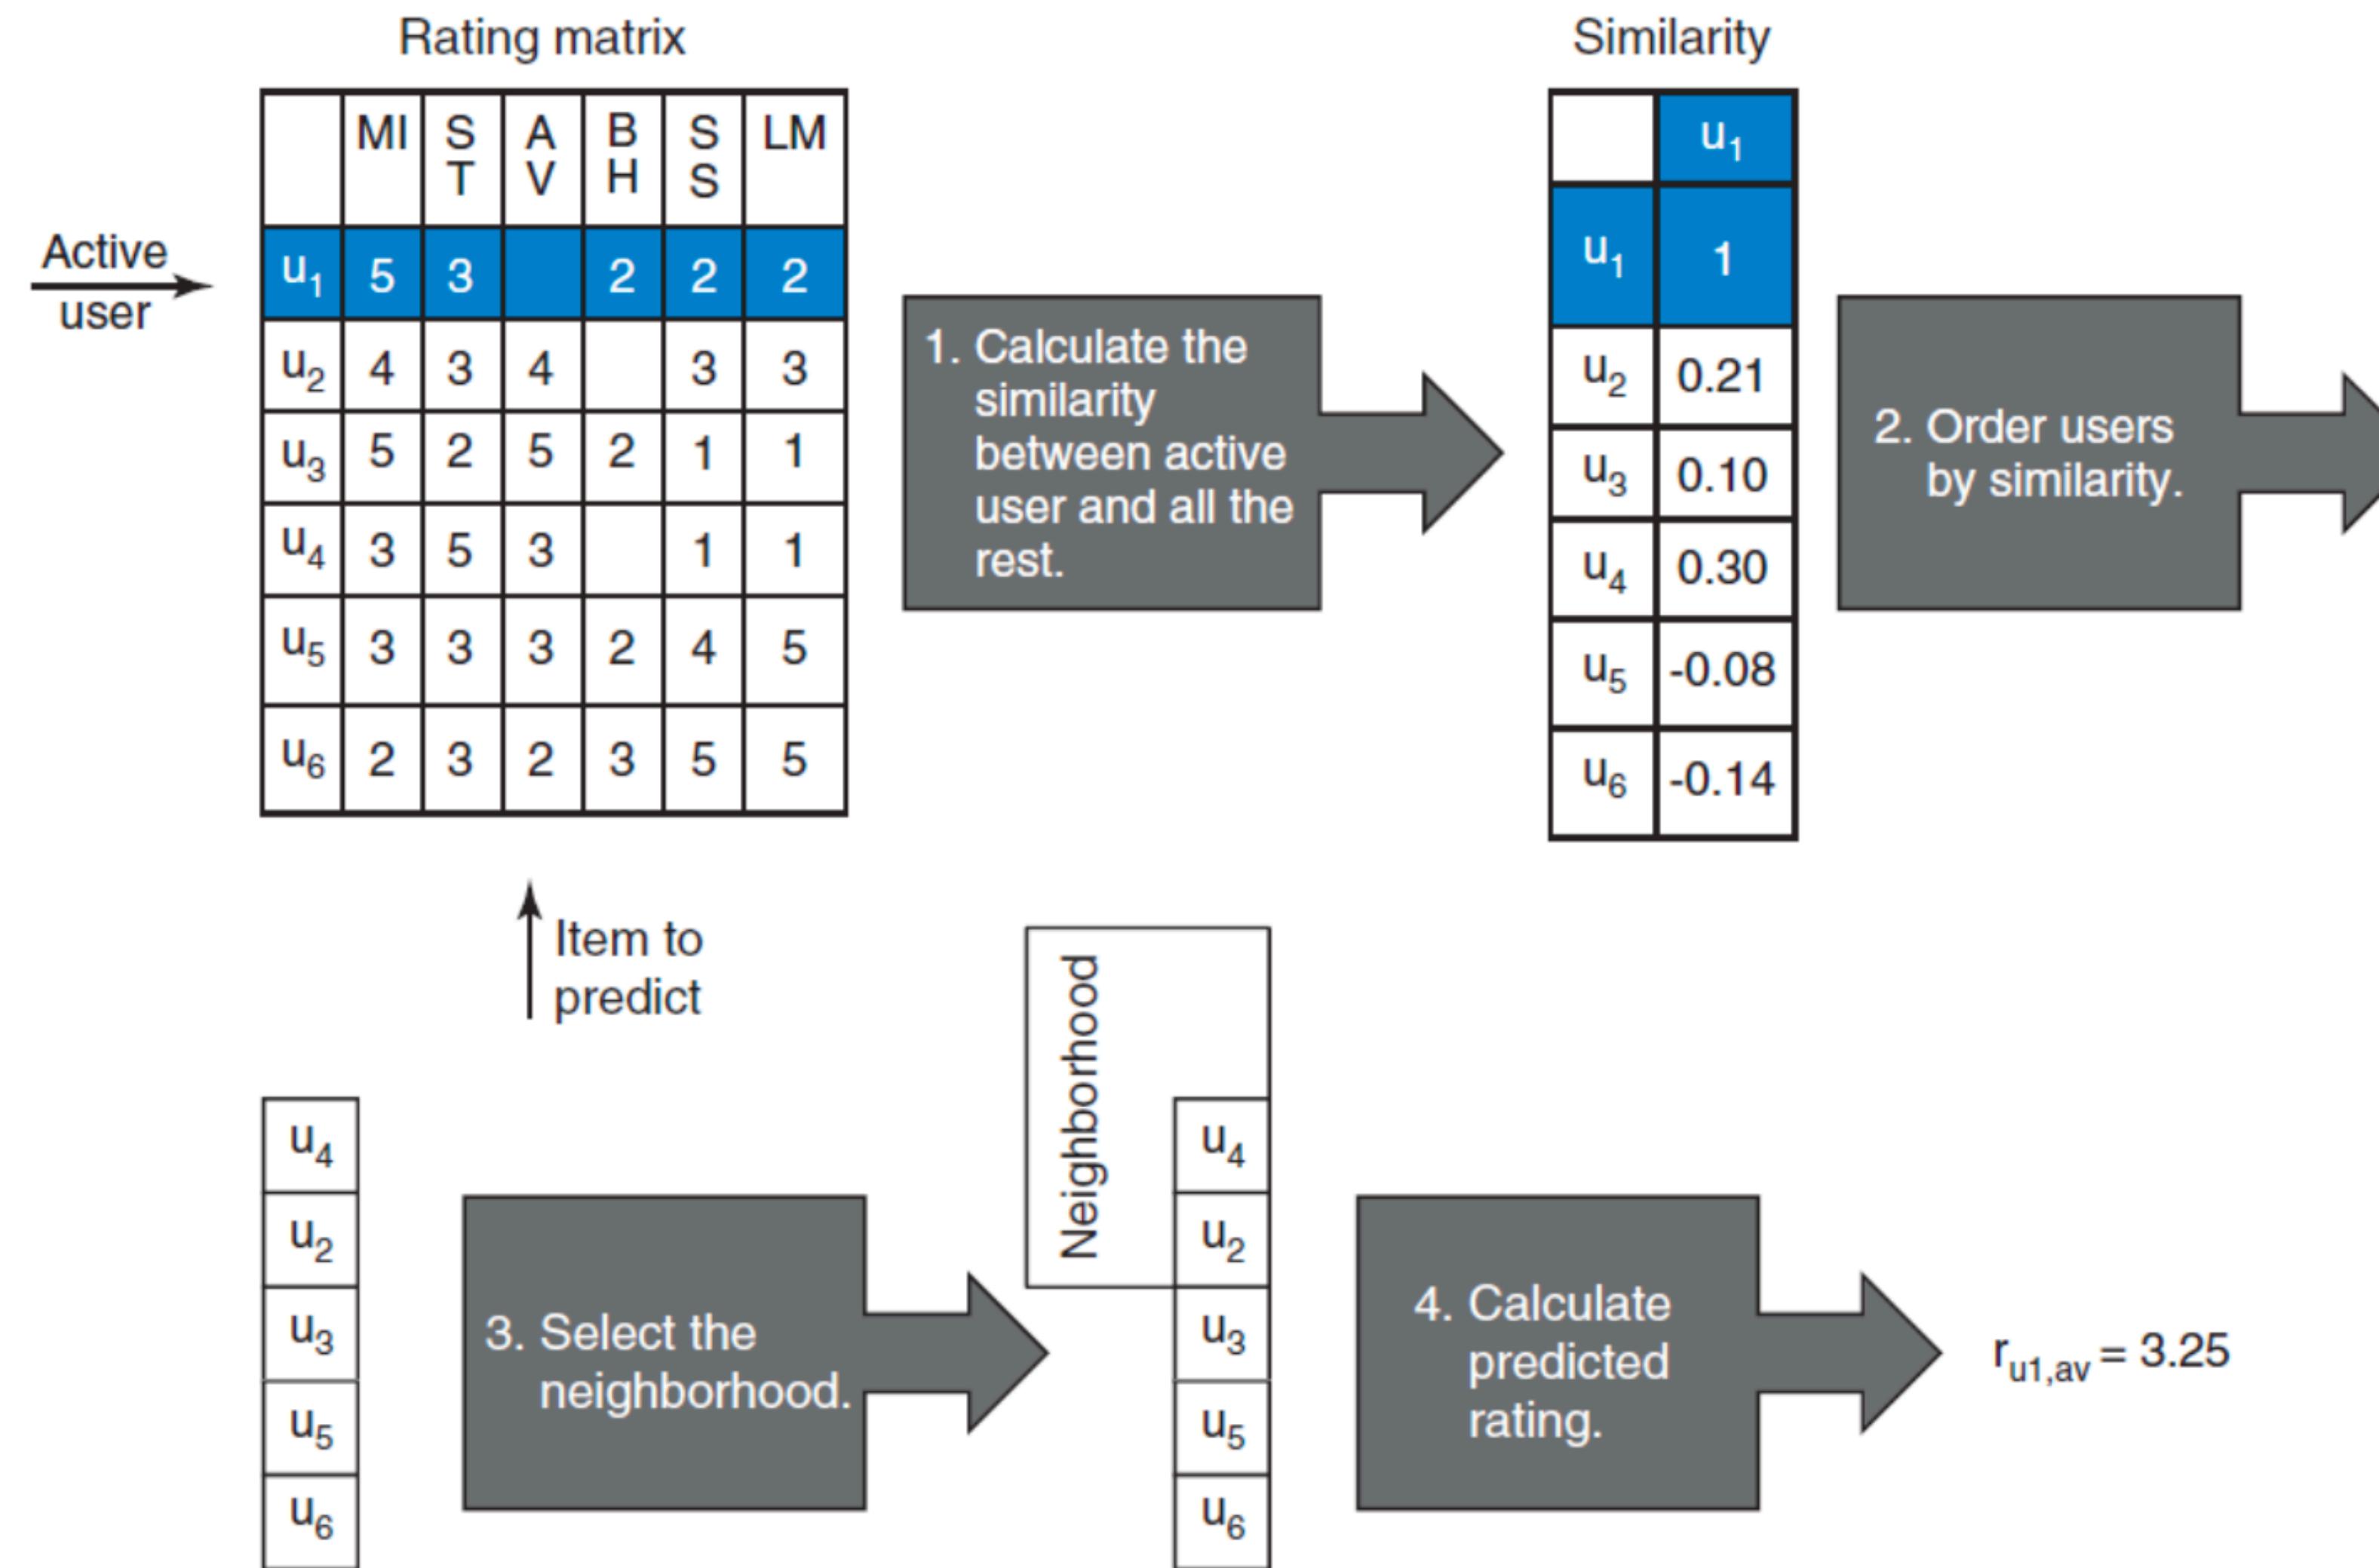
User-based CF

Steps:

1. Determine the similarity of users and a give user (i.e., target user)
2. Aggregate the ratings of similar users to predict target user's ratings
3. Decide how many similar users' opinions should be considered



User-based CF



User-based CF

Step 1: determine the similarity

- Jaccard Distance: $sim(u_i, u_j) = \frac{U_i \cap U_j}{U_i \cup U_j}$
- Cosine Distance: $sim(u_i, u_j) = \frac{\sum_{i=1}^n U_i U_j}{\sqrt{\sum_{i=1}^n U_i^2} \sqrt{\sum_{i=1}^n U_j^2}}$
- Pearson Correlation Coefficient:

$$sim(u_i, u_j) = \frac{\sum_{p \in P} (r_{i,p} - \bar{r}_i)(r_{j,p} - \bar{r}_j)}{\sqrt{\sum_{p \in P} (r_{i,p} - \bar{r}_i)^2} \sqrt{\sum_{p \in P} (r_{j,p} - \bar{r}_j)^2}}$$

Where u_i, u_j is a pair of users; U_i, U_j is item list rated by user u_i, u_j respectively; P is the set of items rated by users u_i and u_j ; \bar{r} is the user's average ratings

User-based CF

Step 1: determine the similarity

	Grey's Anatomy	American Horror Story	Money Heist	Sherlock	Breaking Bad
Lucas	5	3	4	4	?
User 1	3	1	2	3	3
User 2	4	3	4	3	5
User 3	3	3	1	5	4
User 4	1	5	5	2	1

■ Pearson Correlation Coefficient (PCC)

$$= \frac{\sum_{p \in P} (r_{Lucas,p} - \bar{r}_{Lucas})(r_{1,p} - \bar{r}_1)}{\sqrt{\sum_{p \in P} (r_{Lucas,p} - \bar{r}_{Lucas})^2} \sqrt{\sum_{p \in P} (r_{1,p} - \bar{r}_1)^2}} = \frac{(5-4)(3-2.25) + (3-4)(1-2.25) + (4-4)(2-2.25) + (4-4)(3-2.25)}{\sqrt{(5-4)^2 + (3-4)^2 + (4-4)^2 + (4-4)^2} \sqrt{(3-2.25)^2 + (1-2.25)^2 + (2-2.25)^2 + (3-2.25)^2}}$$

Where $\bar{r}_{Lucas} = 4$, and $\bar{r}_1 = 2.25$ and PCC is $\in [-1,1]$

User-based CF

Step 1: determine the similarity

Experience:

- Generate user-user similarity matrix

	Lucas	User 1	User 2	User 3	User 4
Lucas	1				
User 1	X	1			
User 2	X	X	1		
User 3	X	X	X	1	
User 4	X	X	X	X	1

- Empirical analyses show that for user-based CF — and at least for the best studied recommendation domains — the PCC outperforms other measures of comparing users (e.g., Spearman's rank correlation coefficient)

User-based CF

Step 2: aggregate the ratings of similar users to predict

- Calculate if the neighbours' ratings for the unseen item (e.g., Breaking Bad) are higher or lower than their average
- Combine the rating differences – use the similarity as the weight
- Add/subtract the neighbours' bias from the target user's (e.g., Lucas's) average and use this as a prediction
- Different types of aggregation functions:

$$r_{u,i} = k \sum_{u' \in N} sim(u, u') \times r_{u',i}$$

$$r_{u,i} = \bar{r}_u + \frac{\sum_{u' \in N} sim(u, u') \times (r_{u',i} - \bar{r}_{u'})}{\sum_{u' \in N} |sim(u, u')|}$$

User-based CF

Step 3: which and how many neighbours for aggregation

- Only consider positively correlated neighbours (or set a threshold)
- Can be optimized based on dataset
- Experience: empirical studies show that aggregating between 50-200 neighbours could be good enough

User-based CF

Discussion

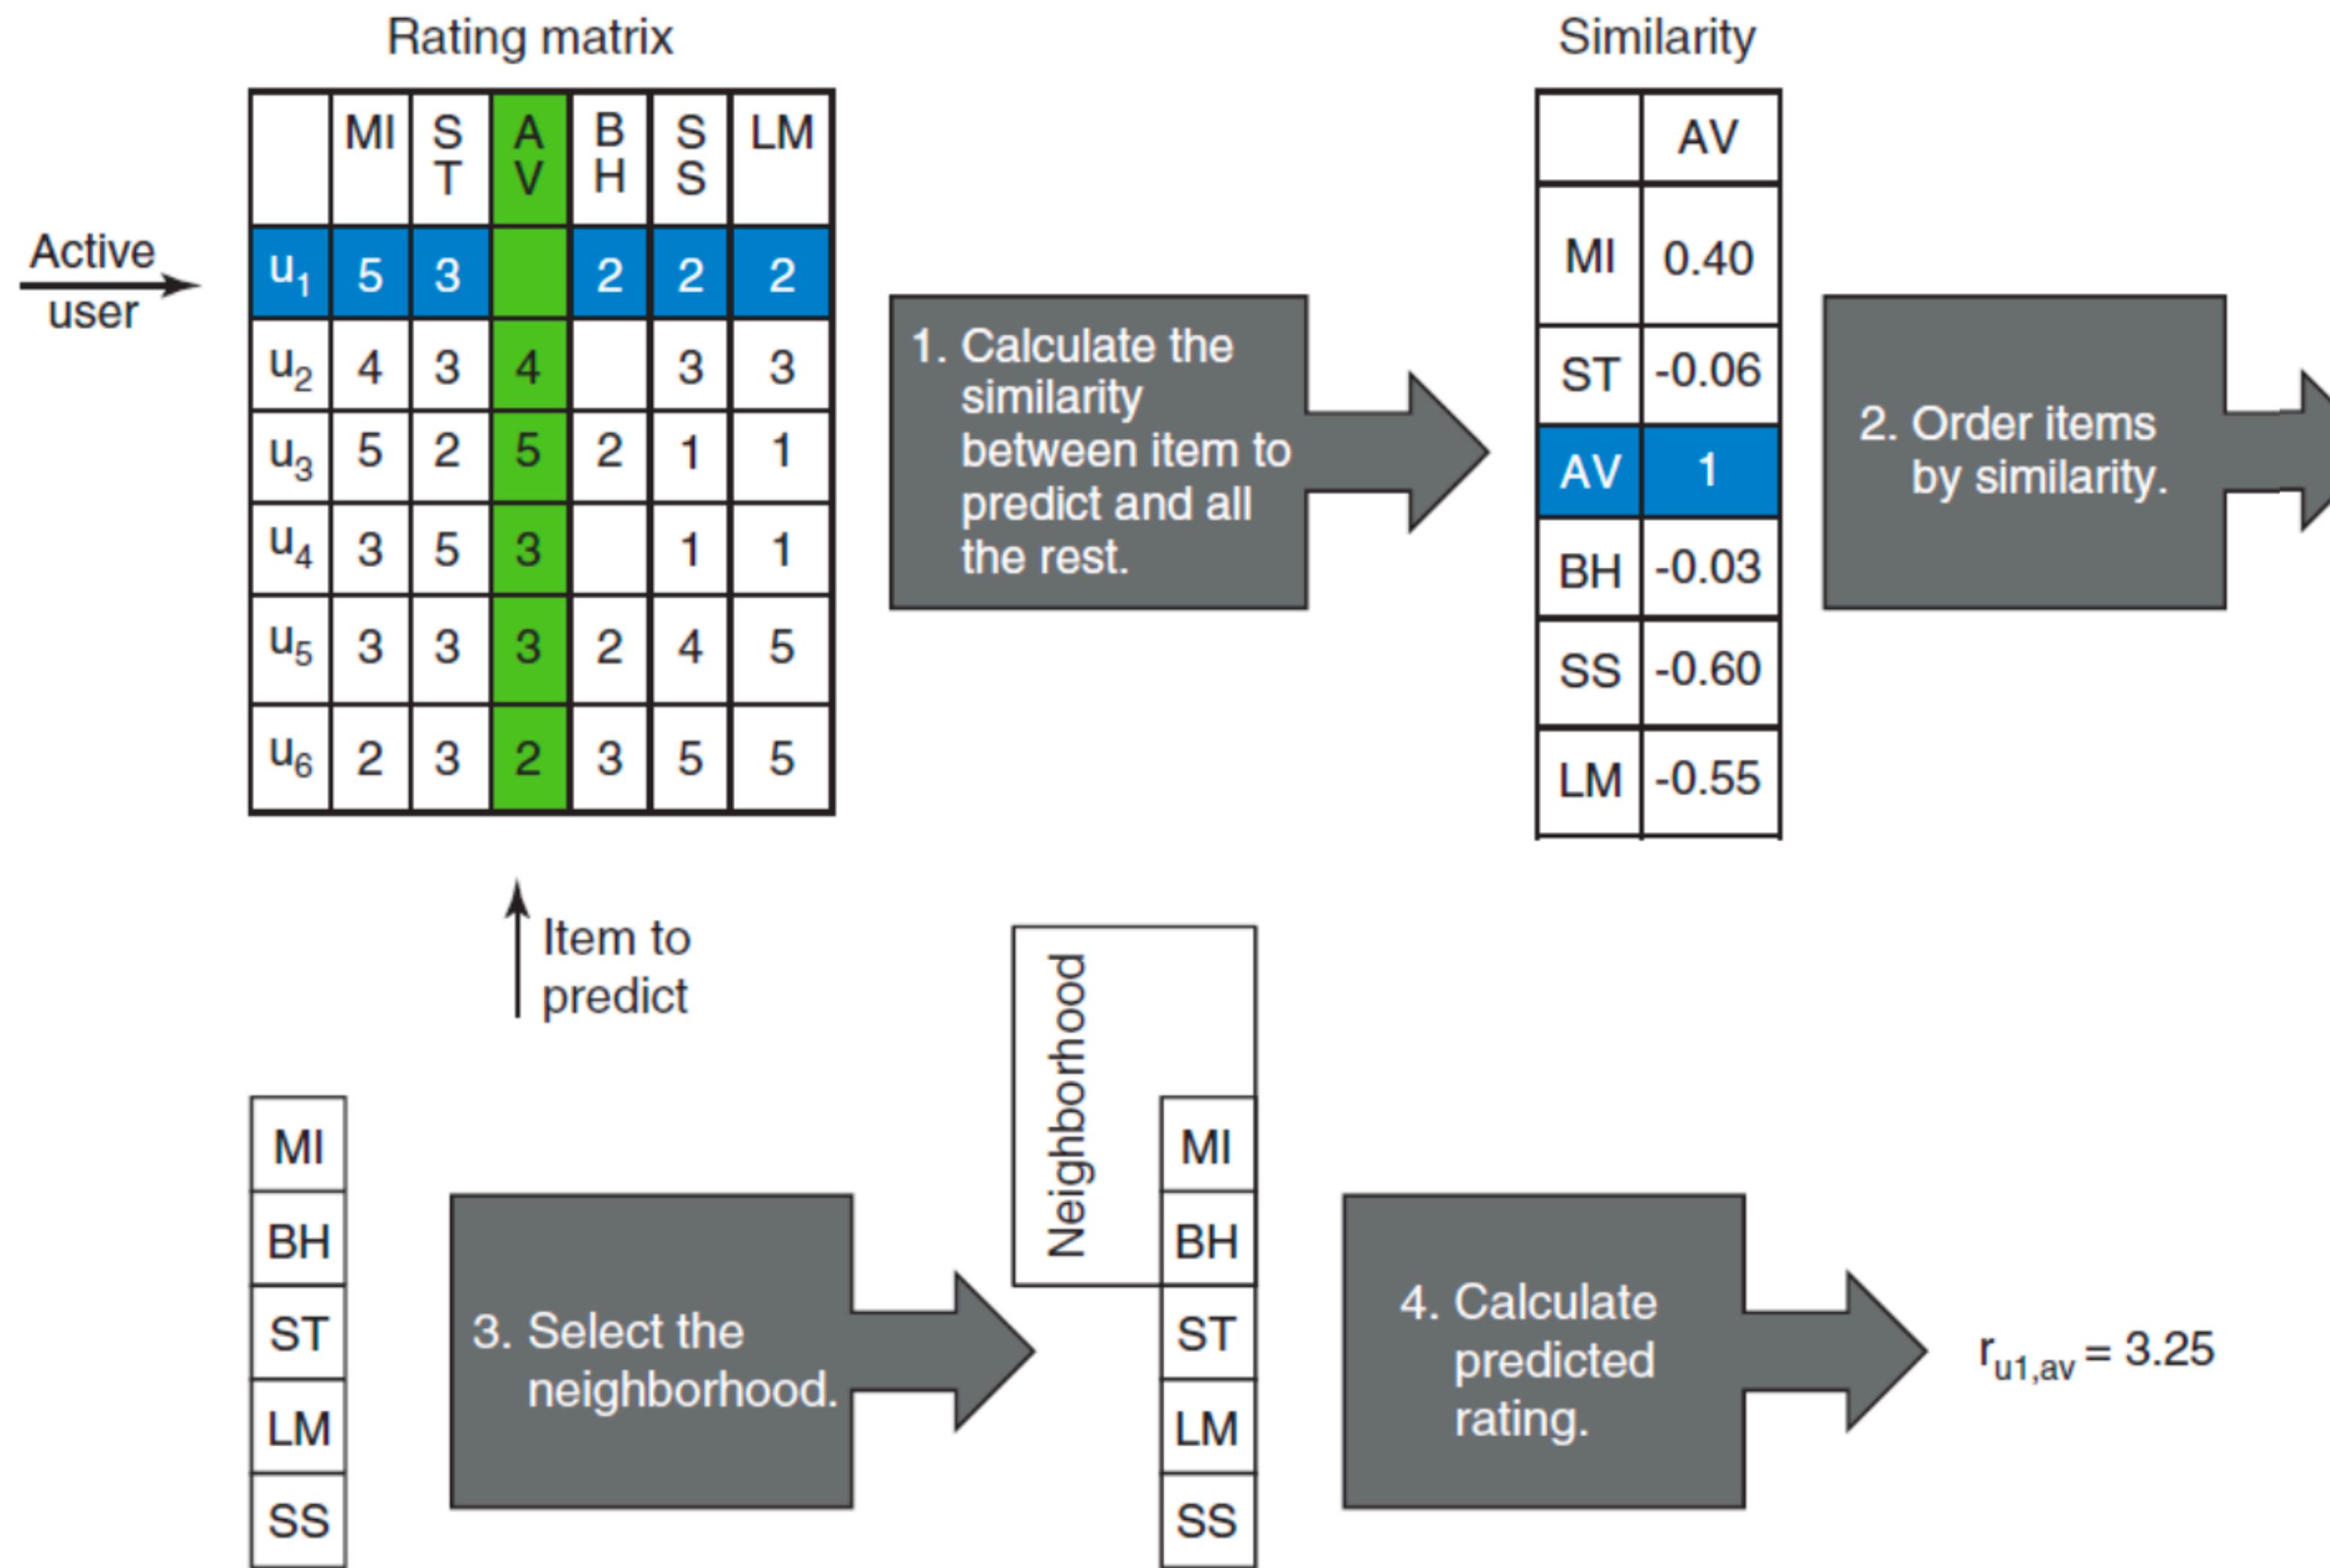
- Very simple scheme leading to quite accurate recommendations
- Possible Issues?

Break

Item-based CF

- Instead of calculating similarity between users, what if we use the similarity between items to make predictions
- Why?
 - The need to scan a vast number of potential neighbours makes it impossible to compute predictions in real-time
 - Item-based is more adaptable for offline pre-processing and thus allows for the computation of recommendations in real-time even for a very large rating matrix

Item-based CF



Item-based CF

Steps:

1. Look into the items that have been rated and compute how similar they are to the target item
2. Select k most similar items
3. Compute prediction by taking weighted average on the target user's ratings on the most similar items



Item-based CF

Step 1: determine the similarity

	Grey's Anatomy	American Horror Story	Money Heist	Sherlock	Breaking Bad
Lucas	5	3	4	4	?
User 1	3	1	2	3	3
User 2	4	3	4	3	5
User 3	3	3	1	5	4
User 4	1	5	5	2	1

- We first compare the rating vectors of the other items and look for items that have ratings similar to target item
- Then, we take Lucas's ratings for these items to predict the rating for the target item (i.e., Breaking Bad)

Item-based CF

Step 1: determine the similarity

	Grey's Anatomy	American Horror Story	Money Heist	Sherlock	Breaking Bad
Lucas	5	3	4	4	?
User 1	3	1	2	3	3
User 2	4	3	4	3	5
User 3	3	3	1	5	4
User 4	1	5	5	2	1

- Cosine similarity between Grey's Anatomy and Breaking Bad

$$= \frac{3 \times 3 + 5 \times 4 + 4 \times 3 + 1 \times 1}{\sqrt{3^2 + 5^2 + 4^2 + 1^2} \sqrt{3^2 + 4^2 + 3^2 + 1^2}} = 0.99$$

- Similar to user-based, we want to generate item-item similarity matrix

	Grey's Anatomy	American Horror Story	Money Heist	Sherlock	Breaking Bad
Grey's Anatomy	1				
American Horror Story	X	1			
Money Heist	X	X	1		
Sherlock	X	X	X	1	
Breaking Bad	X	X	X	X	1

Item-based CF

Step 2, 3: aggregate and predict

- After the similarities between the items are determined (i.e., the item-item similarity matrix is generated), we predict the rating of Lucas for Breaking Bad as:

$$r_{u,i} = \frac{\sum_{i' \in N} sim(i, i') \times r_{u,i'}}{\sum_{i' \in N} |sim(i, i')|} = \frac{\sum_{i' \in N} sim(BB, i') \times r_{L,i'}}{\sum_{i' \in N} |sim(BB, i')|}$$

Where N is the set of neighbourhood items

Item-based CF

Discussion

- Bottleneck – similarity computation
- Time complexity: millions of users and items
- Possible Issues?

Item-based CF

Amazon.com Recommendations

- Generally, it's assumed that item similarity is stable, so item similarities can be calculated beforehand or offline
- Generate a data set where you can look up similar items for the current item, making it faster to calculate the predictions.

For each item in product catalog, I_1

For each customer C who purchased I_1

For each item I_2 purchased by
customer C

Record that a customer purchased I_1
and I_2

For each item I_2

Compute the similarity between I_1 and I_2

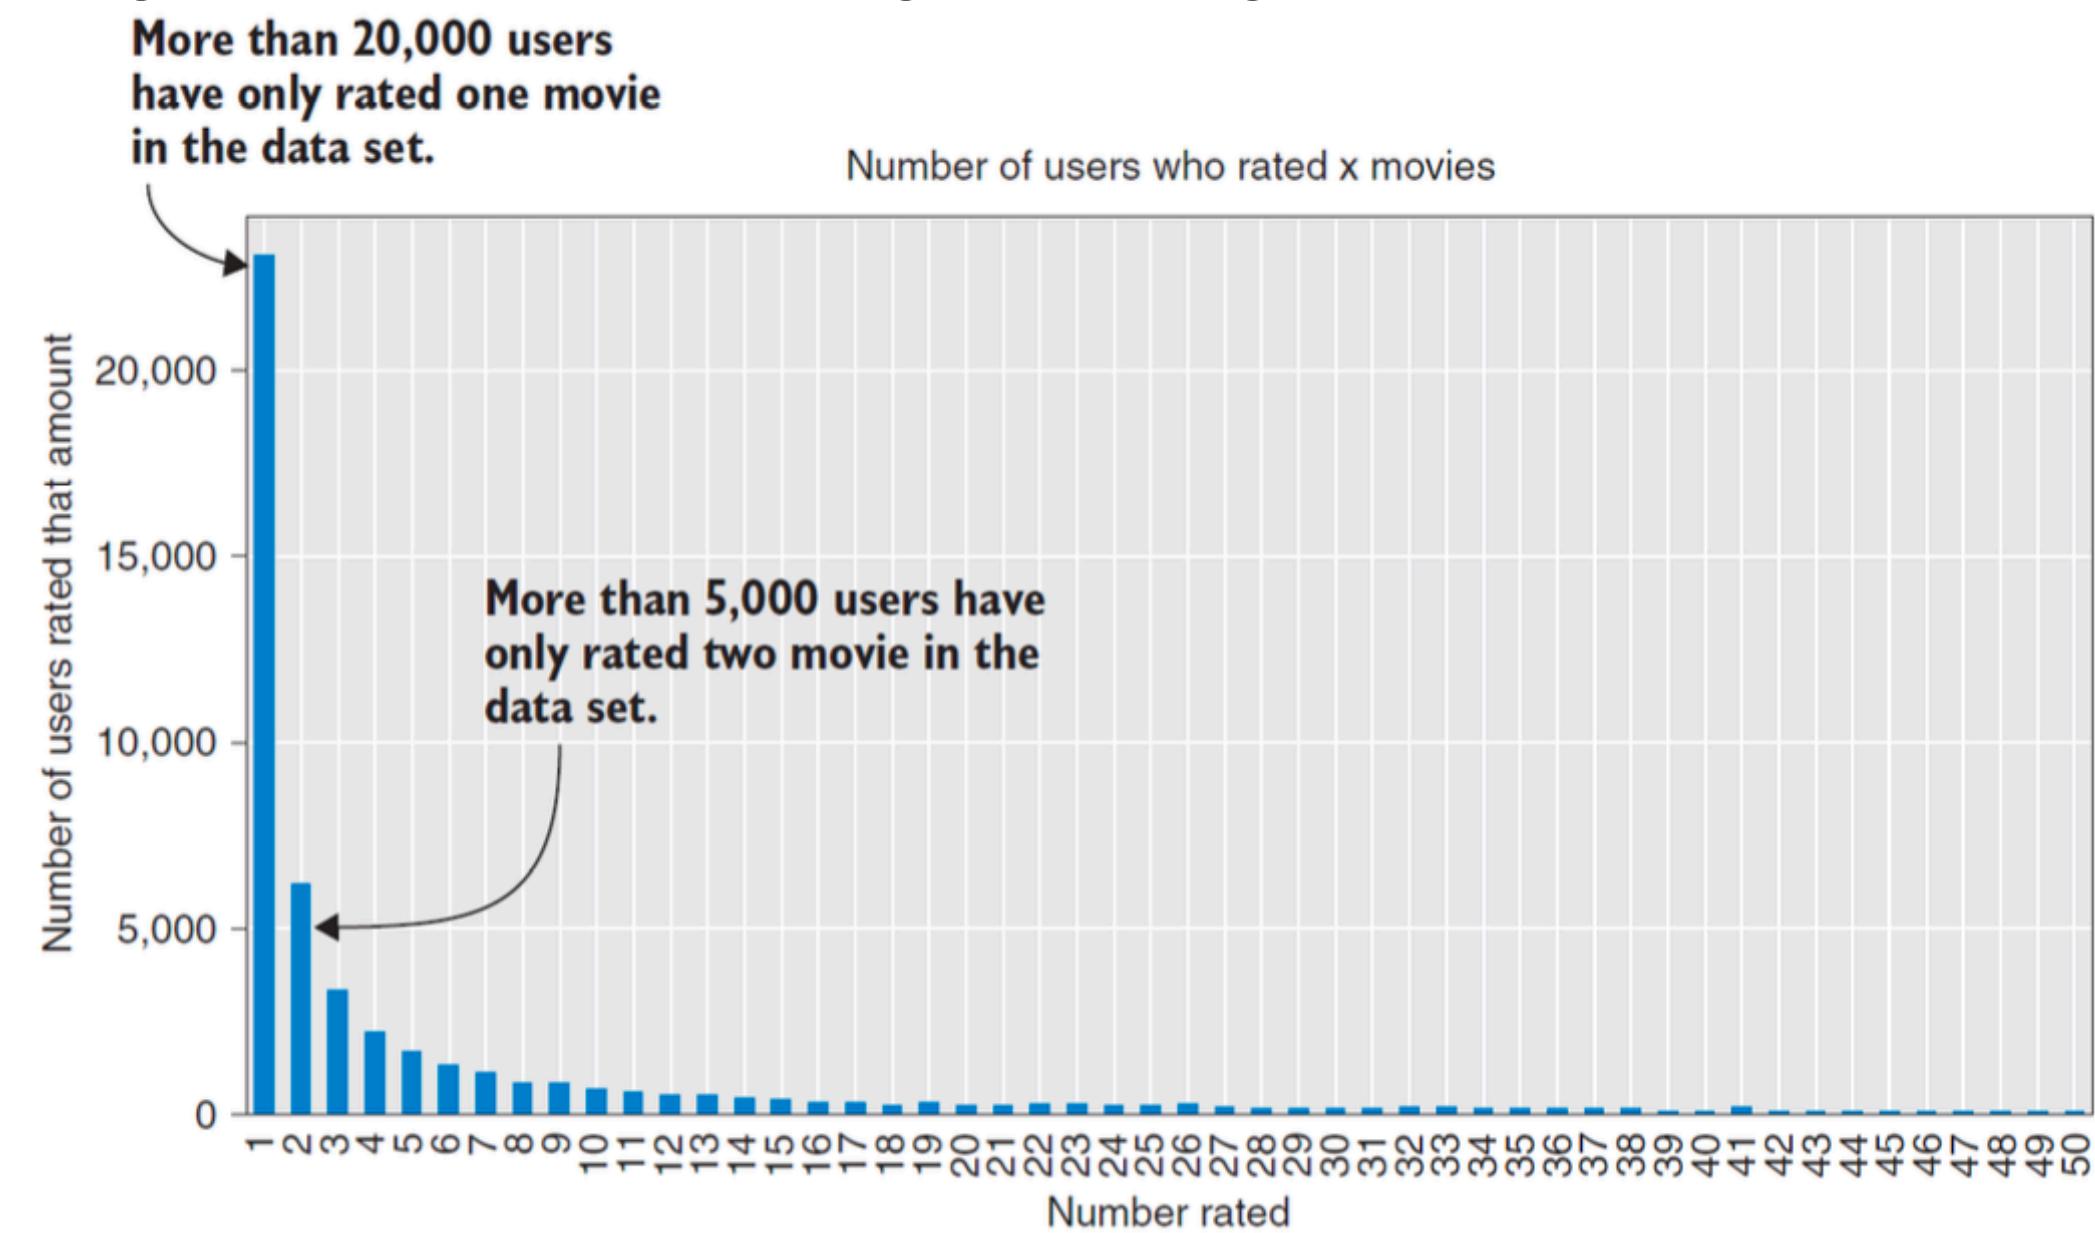
Recap - Neighbourhood-based CF

Recap

- The rating matrix is directly used to find neighbours / make predictions
- User-based CF: similarity between users is dynamic, pre-computing user neighbourhood can lead to poor prediction
- Item-based CF: similarity between items is static, as such pre-computing is possible
- Choice of similarity functions: Pearson Correlation Coefficient (PCC), Cosine, Adjusted Cosine, Euclidean, etc.

Recap

- Some items that only be rated very rarely

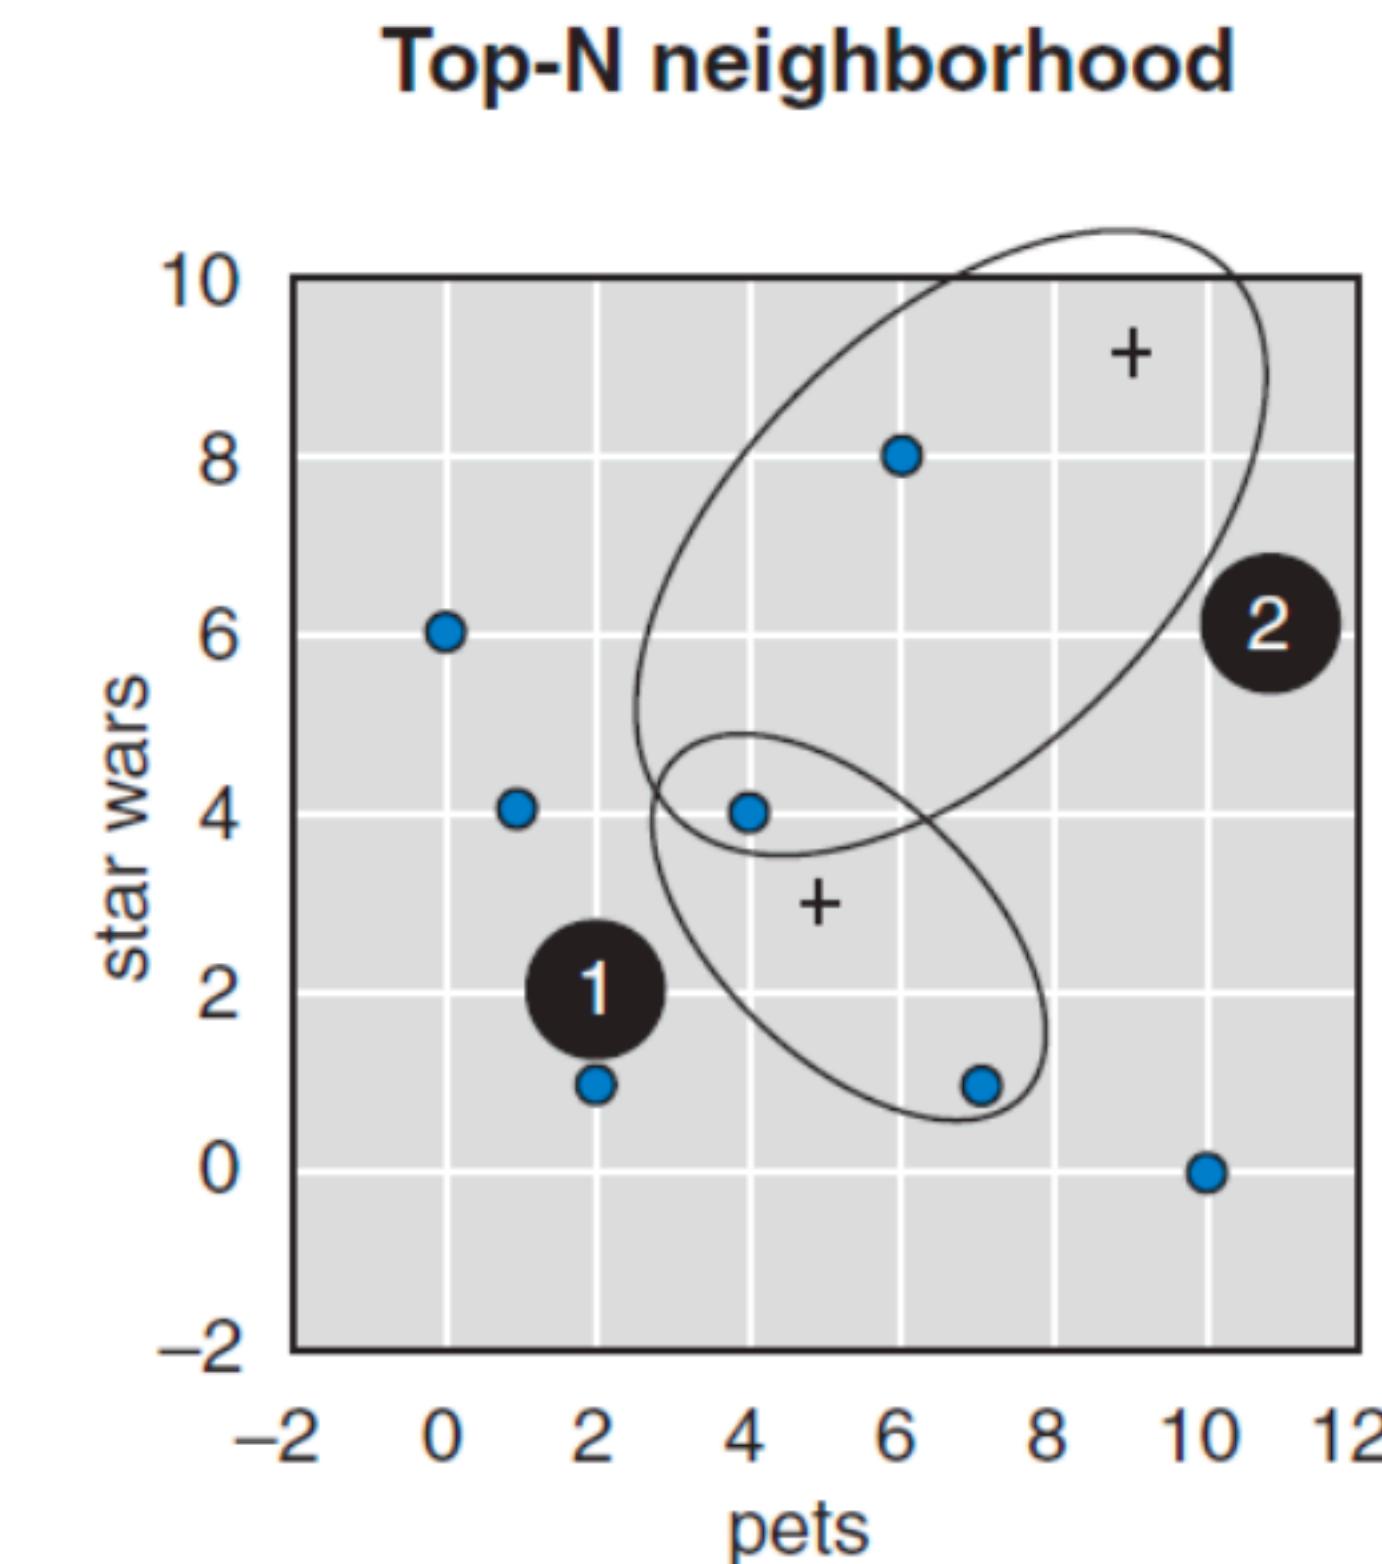
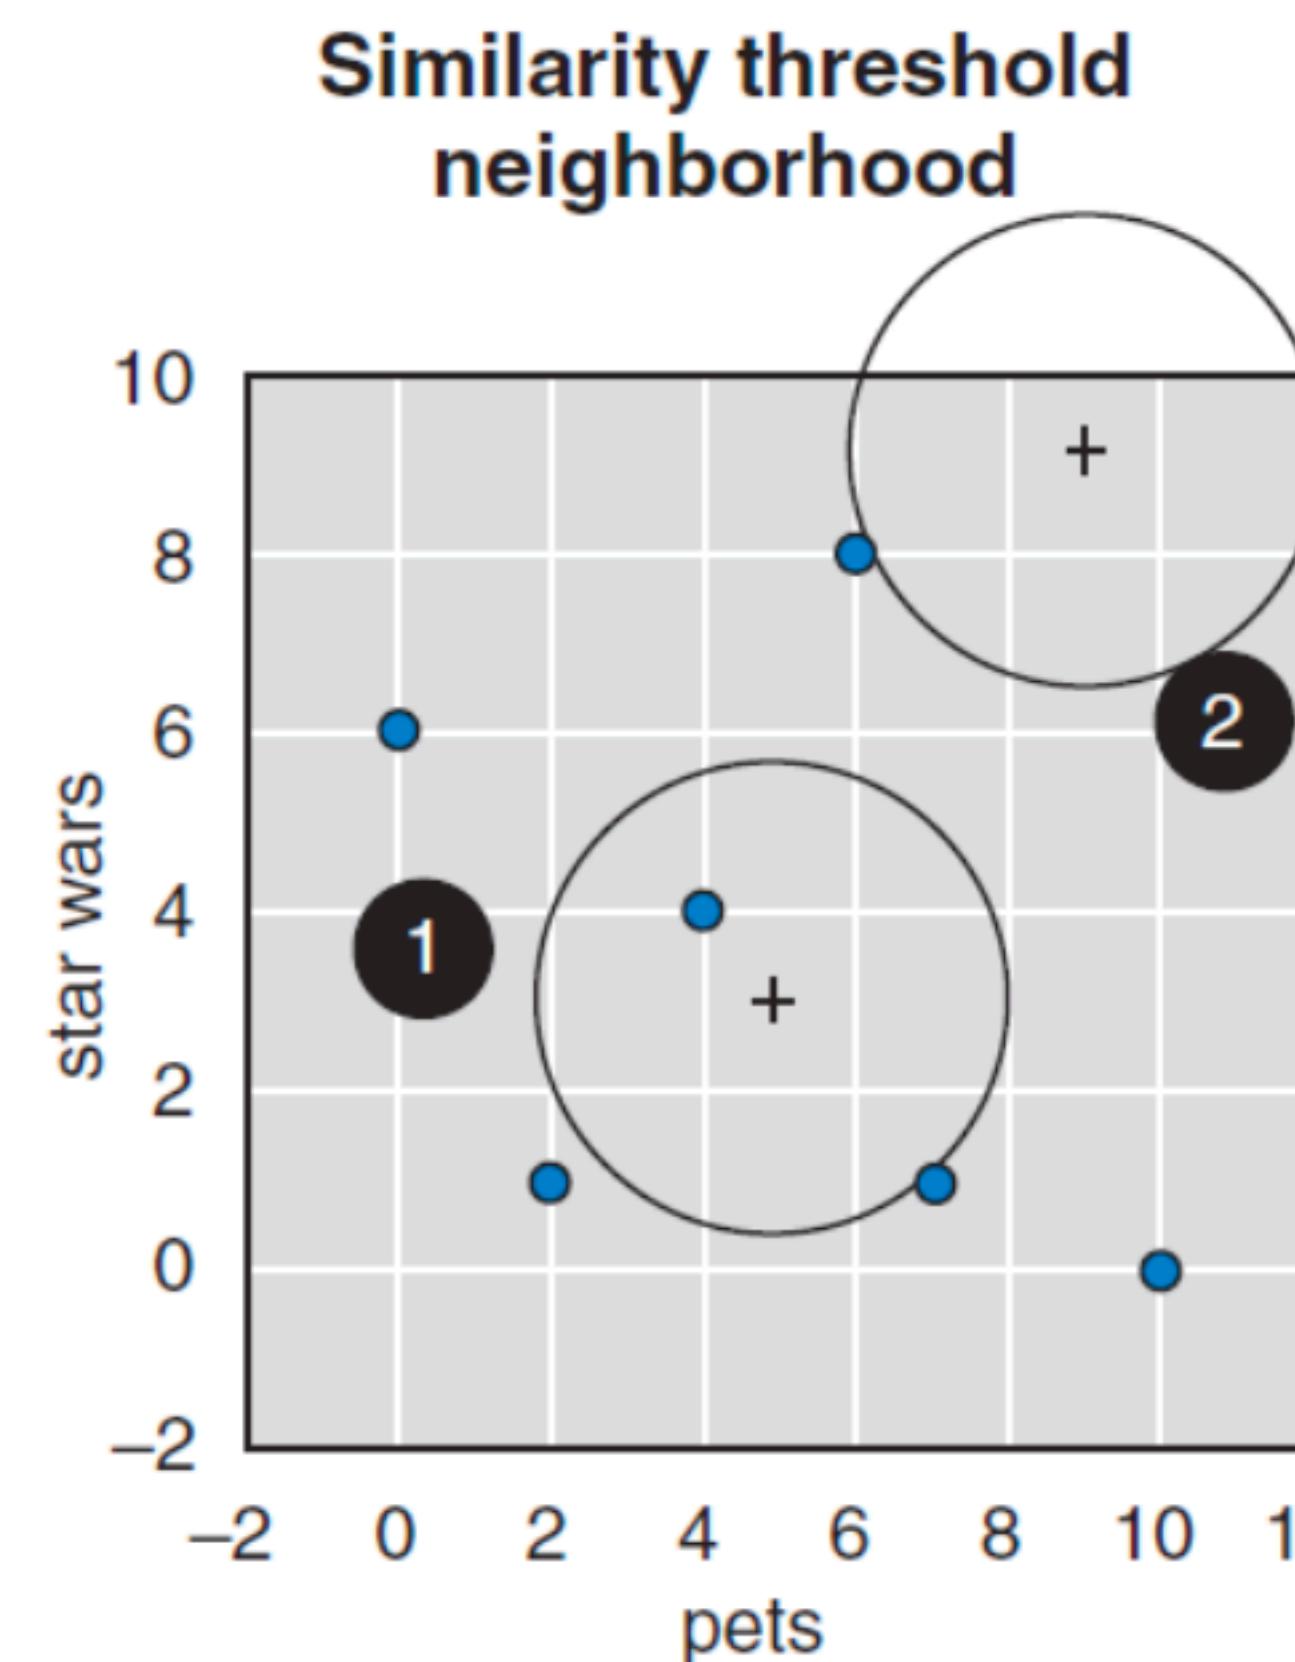


- If an item has only **one** rating, the average rating is equal to itself. Thus, its means is 0, which makes the similarity function undefined
- How about a user who rated **two** items?

Recap

Neighbourhood selection

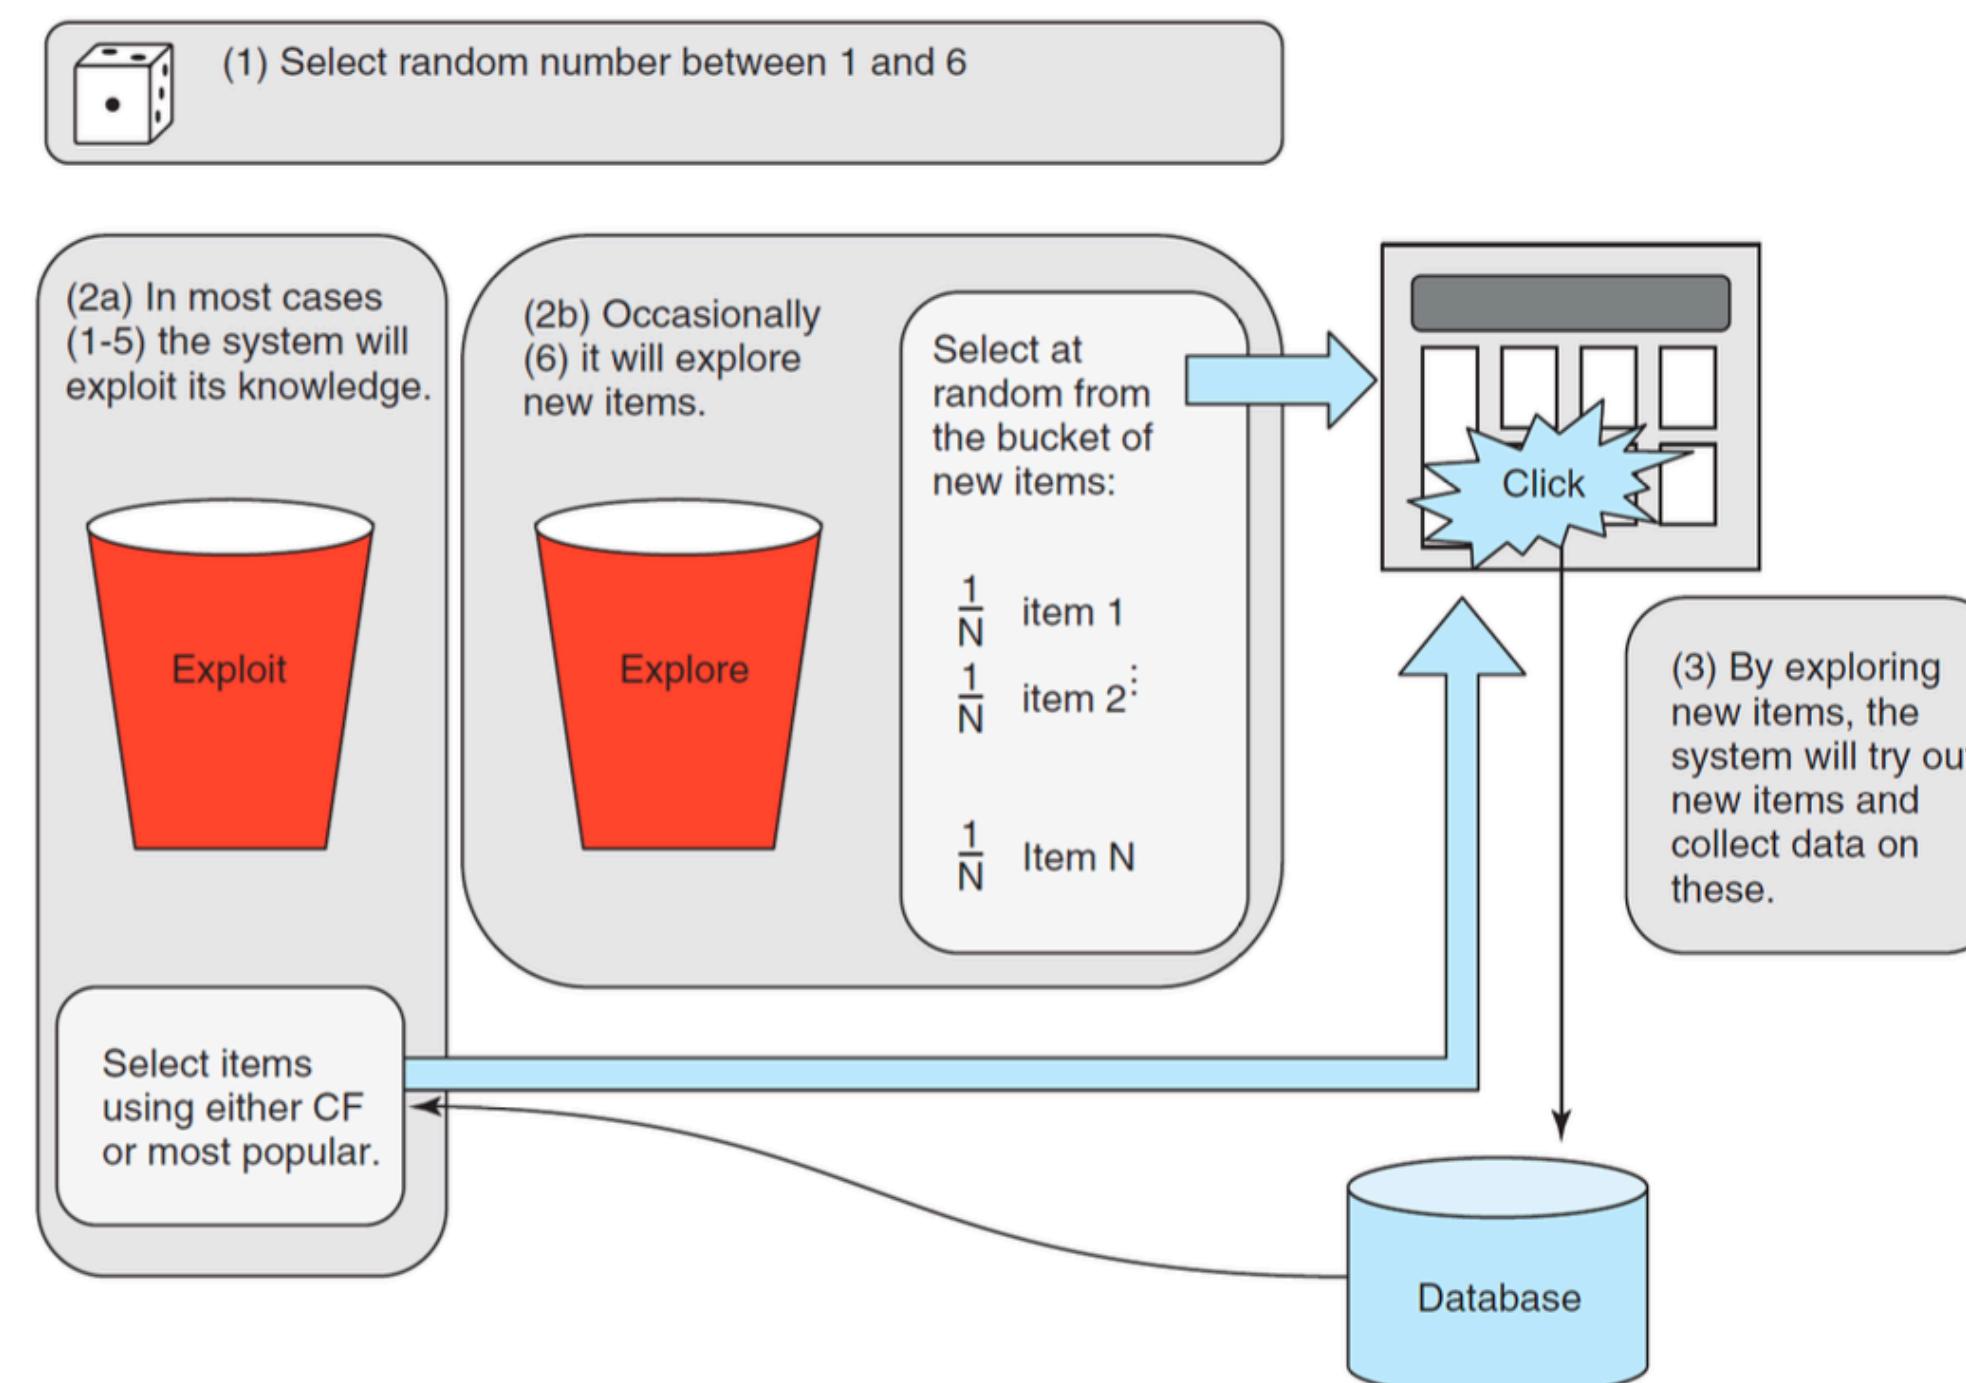
- Clustering, threshold, top-N



Recap

Cold-start

- Explore: show a new item, in which the items gets exposed and users hopefully interact with them
- Exploit: exploit the knowledge you have and return recommendations



Evaluation

Evaluation

In practice

- Different perspective/aspects
- Retrieval perspective
 - Reduce search costs and provide “correct” recommendations
 - Assumption: users already know what they want in advance
- Recommendation perspective
 - Serendipity: identify items from the Long Tail (i.e., users did not know)
 - Provide “appropriate” recommendations

Evaluation

In research

- Offline experiments: historical data, prediction accuracy, etc.
- Test with real users:
 - A/B tests
 - Sales increases/decreases, CTR (click-through rate)
- Simulator

Evaluation

What to evaluate

- Understand users' tastes by minimizing the prediction error
- Coverage: ensure the algorithm will recommend everything in your catalog and whether it can recommend something to all registered users
- Diversity: help users navigate a larger catalog than what is found on a single page
- Serendipity: give users 'sensation', so that all their visits aren't more of the same

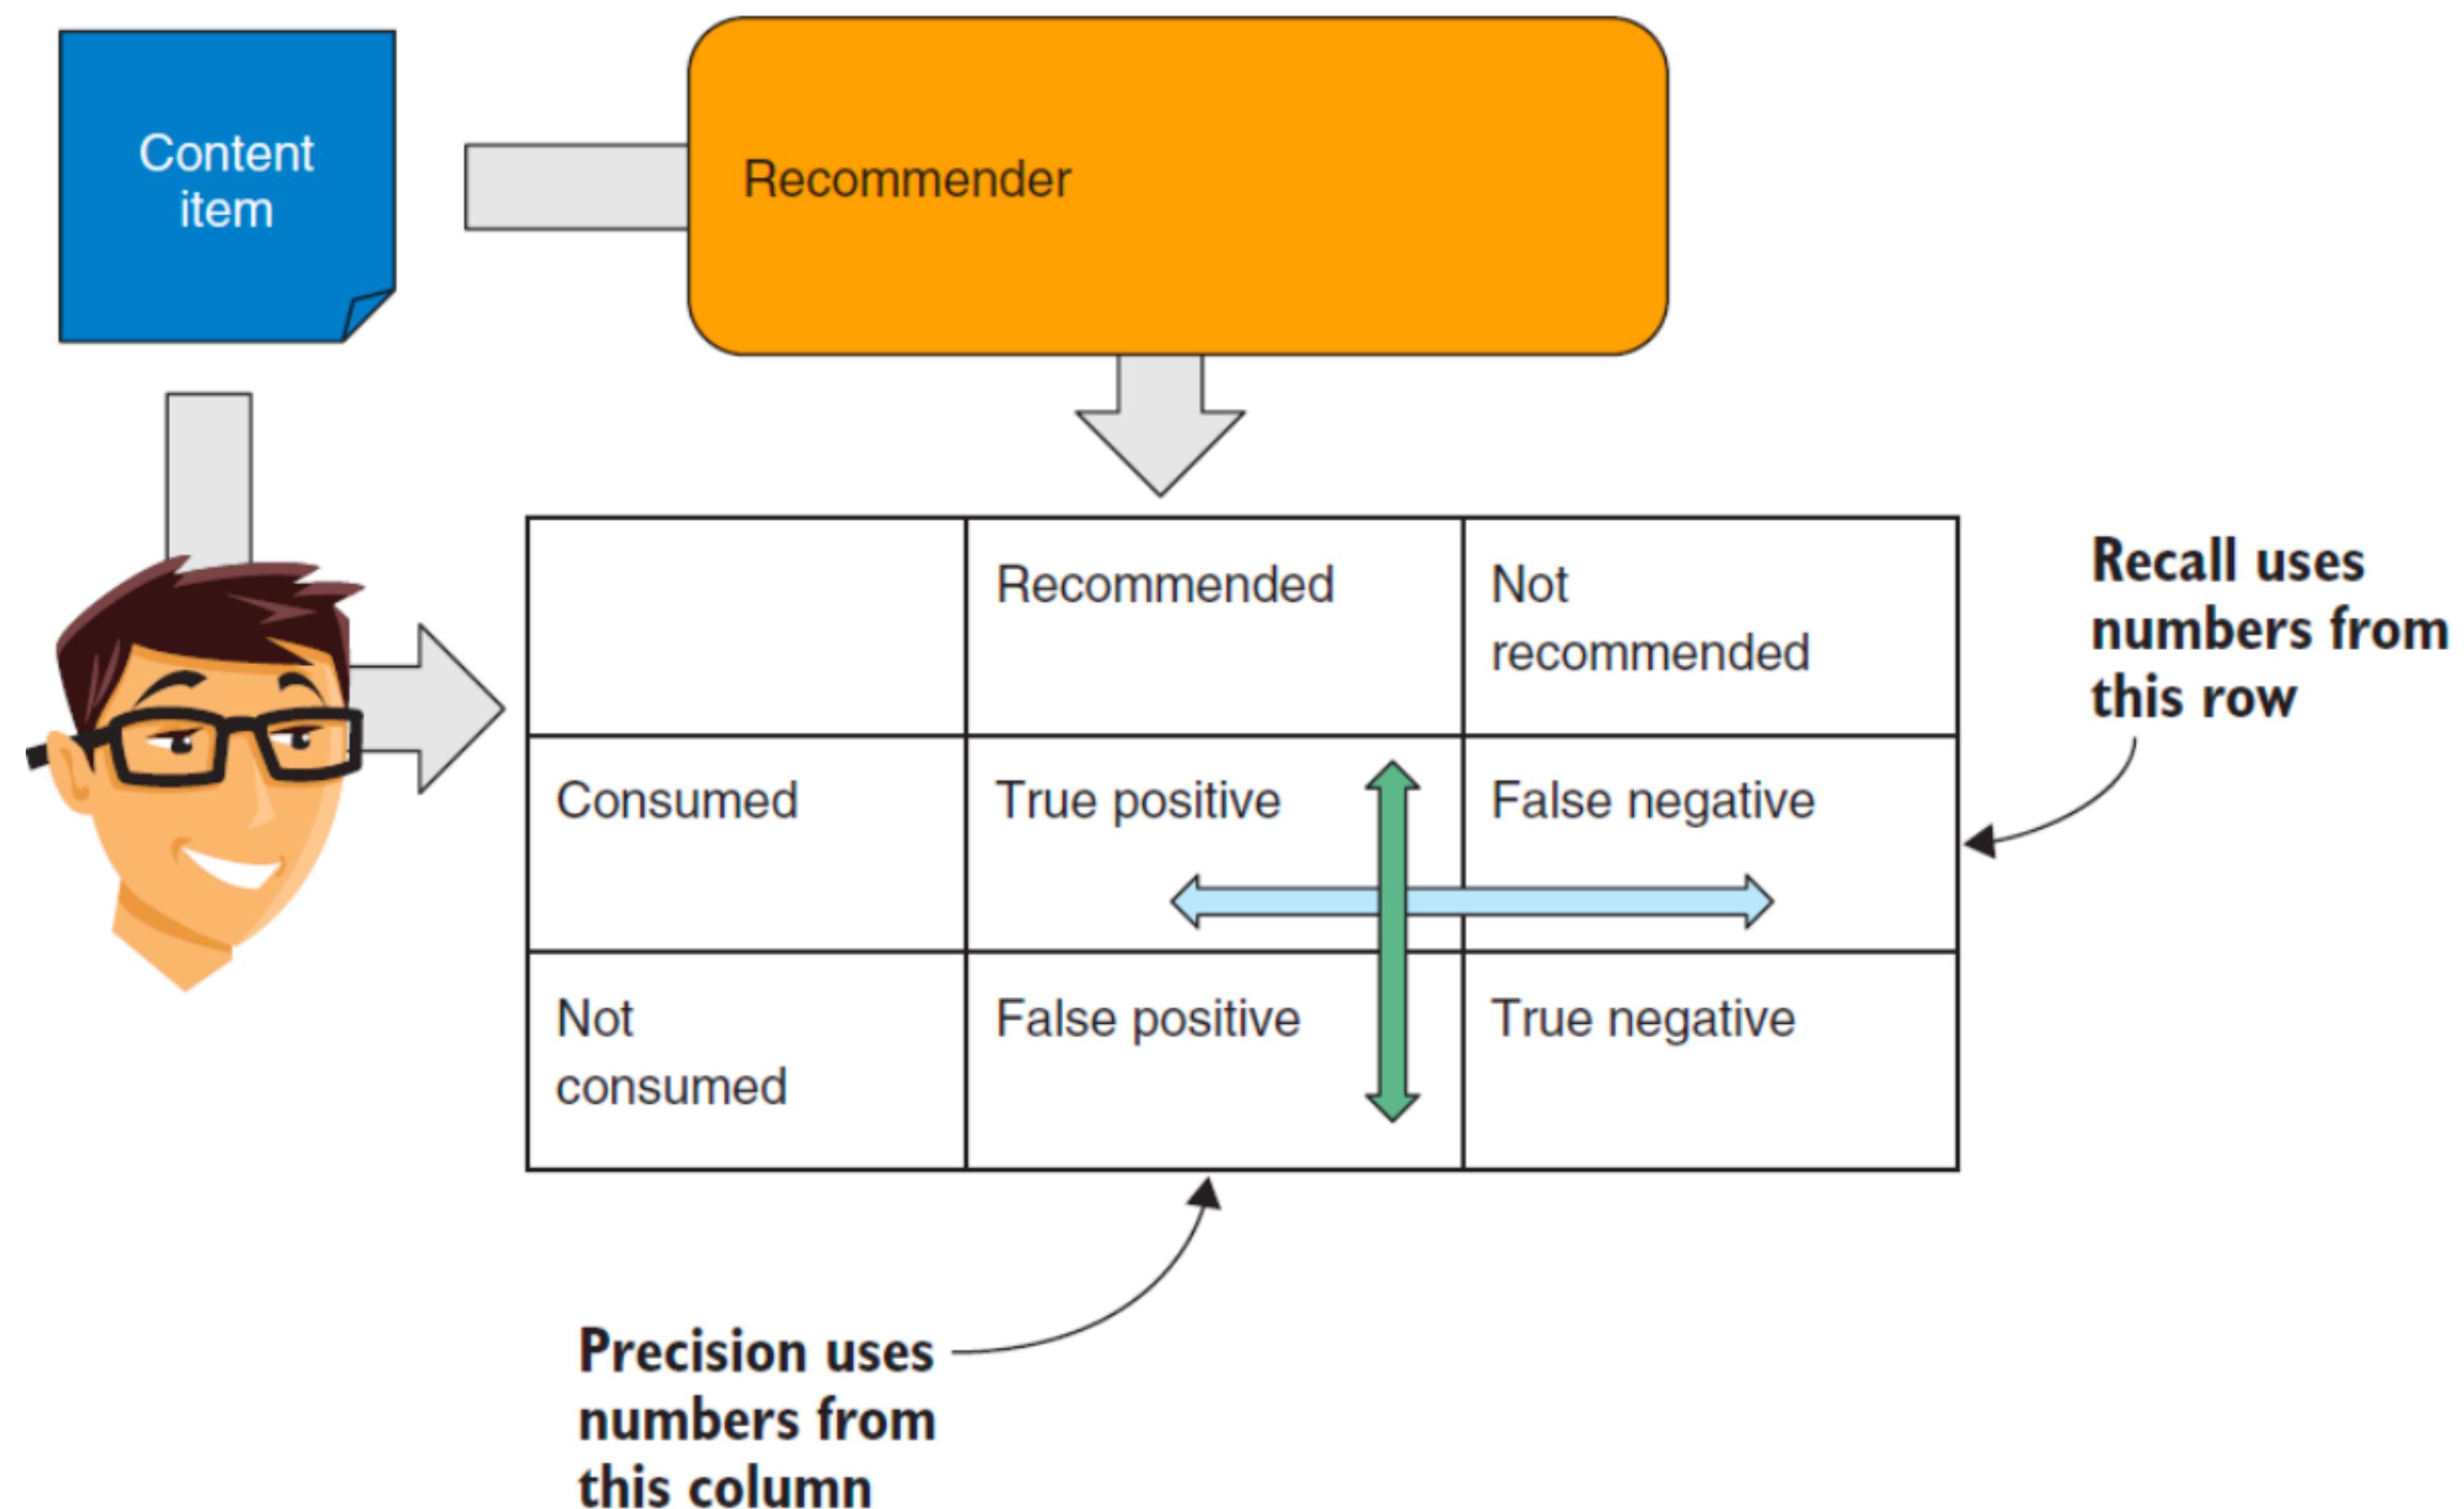
Evaluation Accuracy

- Precision: a measure of exactness, determines the fraction of relevant items retrieved out of all items retrieved
 - E.g., the proportion of recommended movies that are actually good
- Recall: a measure of completeness, determines the fraction of relevant items retrieved out of all relevant items
 - E.g., the proportion of all good movies recommended

$$\text{Precision@K} = \frac{tp}{tp + fp} = \frac{|\text{appropriate movies recommended}|}{|\text{all recommendations}|}$$

$$\text{Recall@K} = \frac{tp}{tp + fn} = \frac{|\text{appropriate movies recommended}|}{|\text{all appropriate movies}|}$$

Evaluation Accuracy



Evaluation

Error rate

- Mean Absolute Error (MAE): computes the deviation between predicted ratings and actual ratings
- Root Mean Square Error (RMSE): similar to MAE, with more emphasis on larger deviation

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |pred_i - r_i| \quad \text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (pred_i - r_i)^2}$$

Evaluation

Rank Score

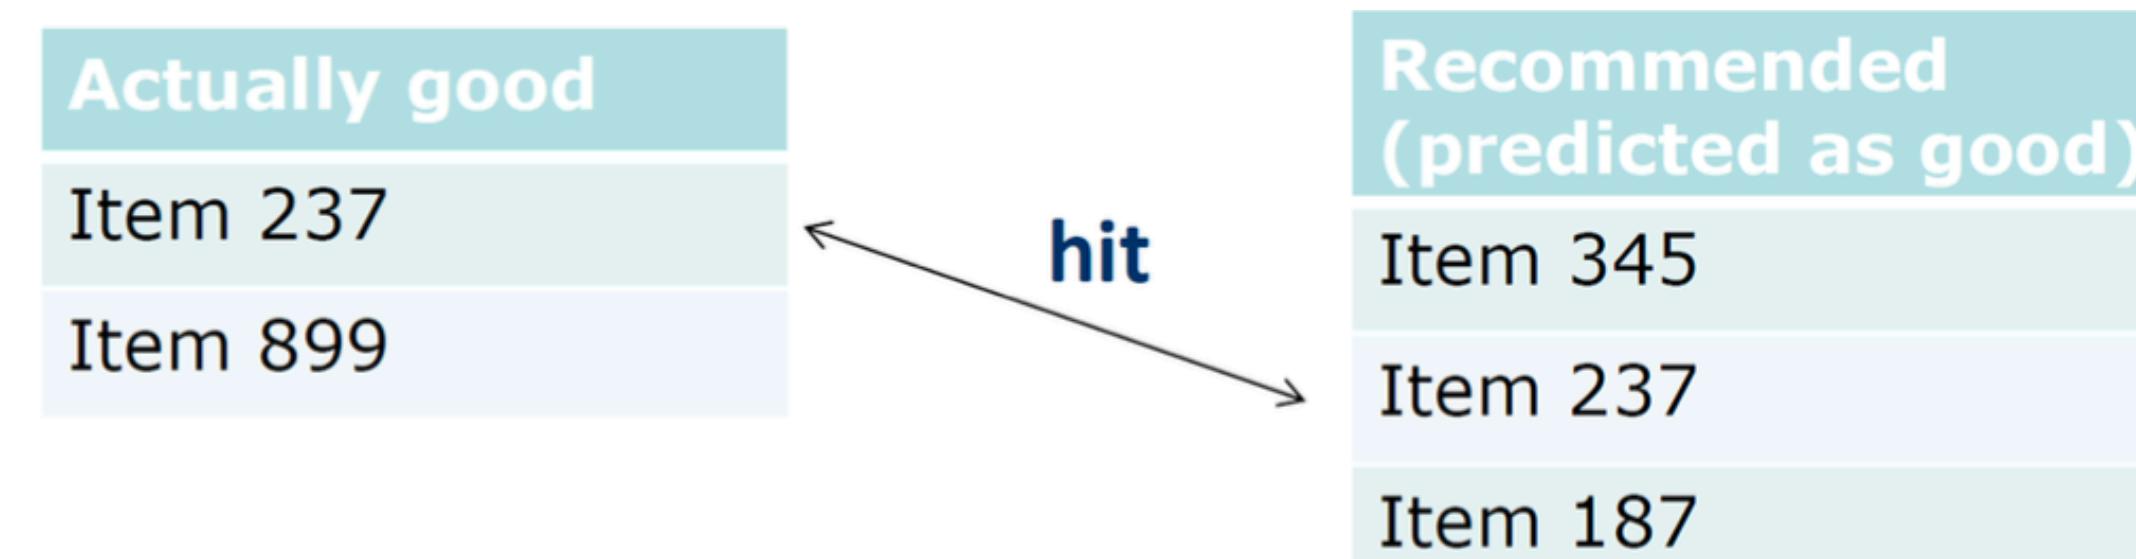
- Extends recall and precision to take the positions of correct items in a ranked list into account
- This is particularly important in recommender systems because lower ranked items may be overlooked by users
- Learning-to-rank: optimise models for such measures (e.g., AUC)



Evaluation

Rank Score

- Hit Ratio: generate top-K recommendations of a target user (i.e., masking rated items from the training set, then use the rest to train to produce top-K recommendations)
- Average reciprocal hit rate (ARHR)
- AUC: measures the likelihood that a random relevant item is ranked higher than a random irrelevant item

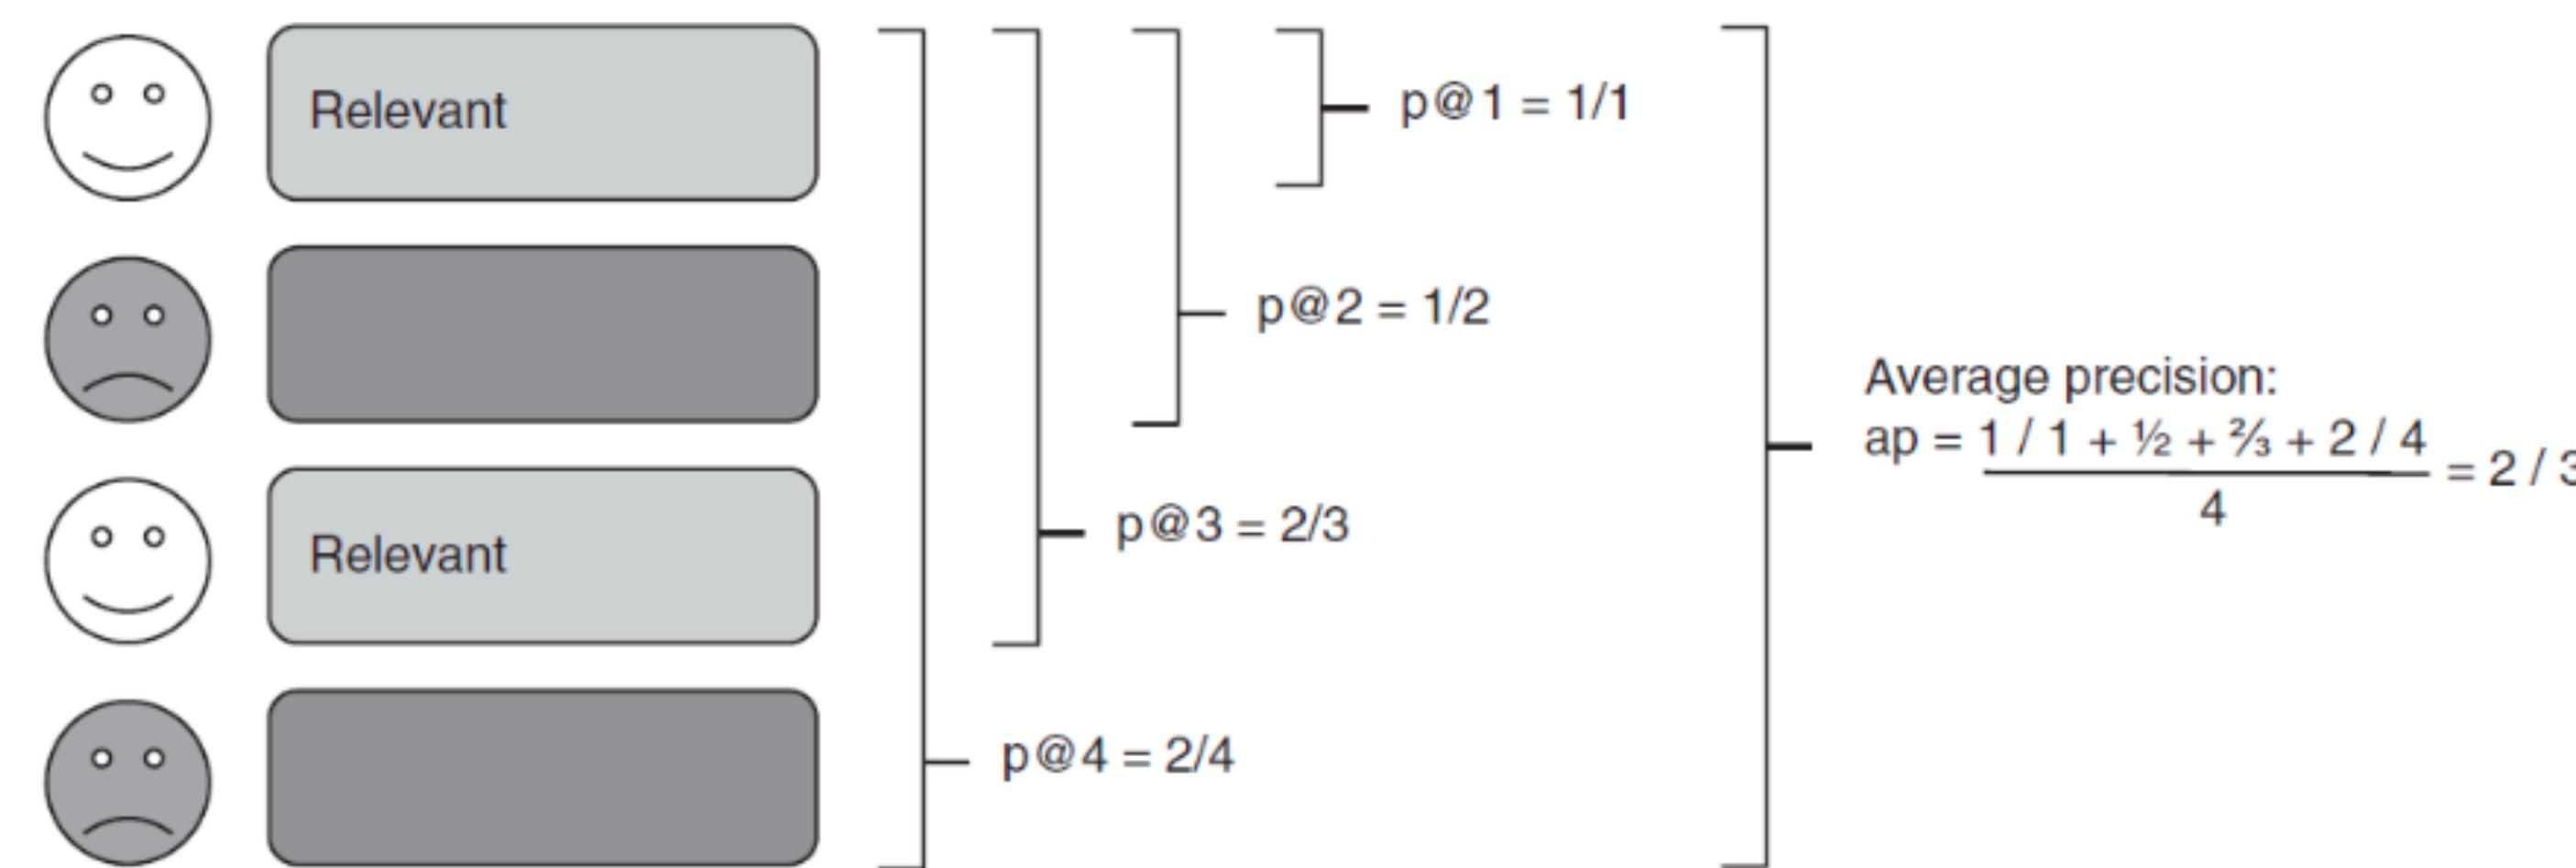


Evaluation

Rank Score

- Mean Average Precision (MAP): measure how good the rank is by running the precision from 1 to m , where m is k items that are recommended

$$AP(U) = \frac{\sum_{k=1}^m Precision @ k(u)}{m}; \quad MAP(U) = \frac{\sum_{u \in U} AP(u)}{|U|}$$



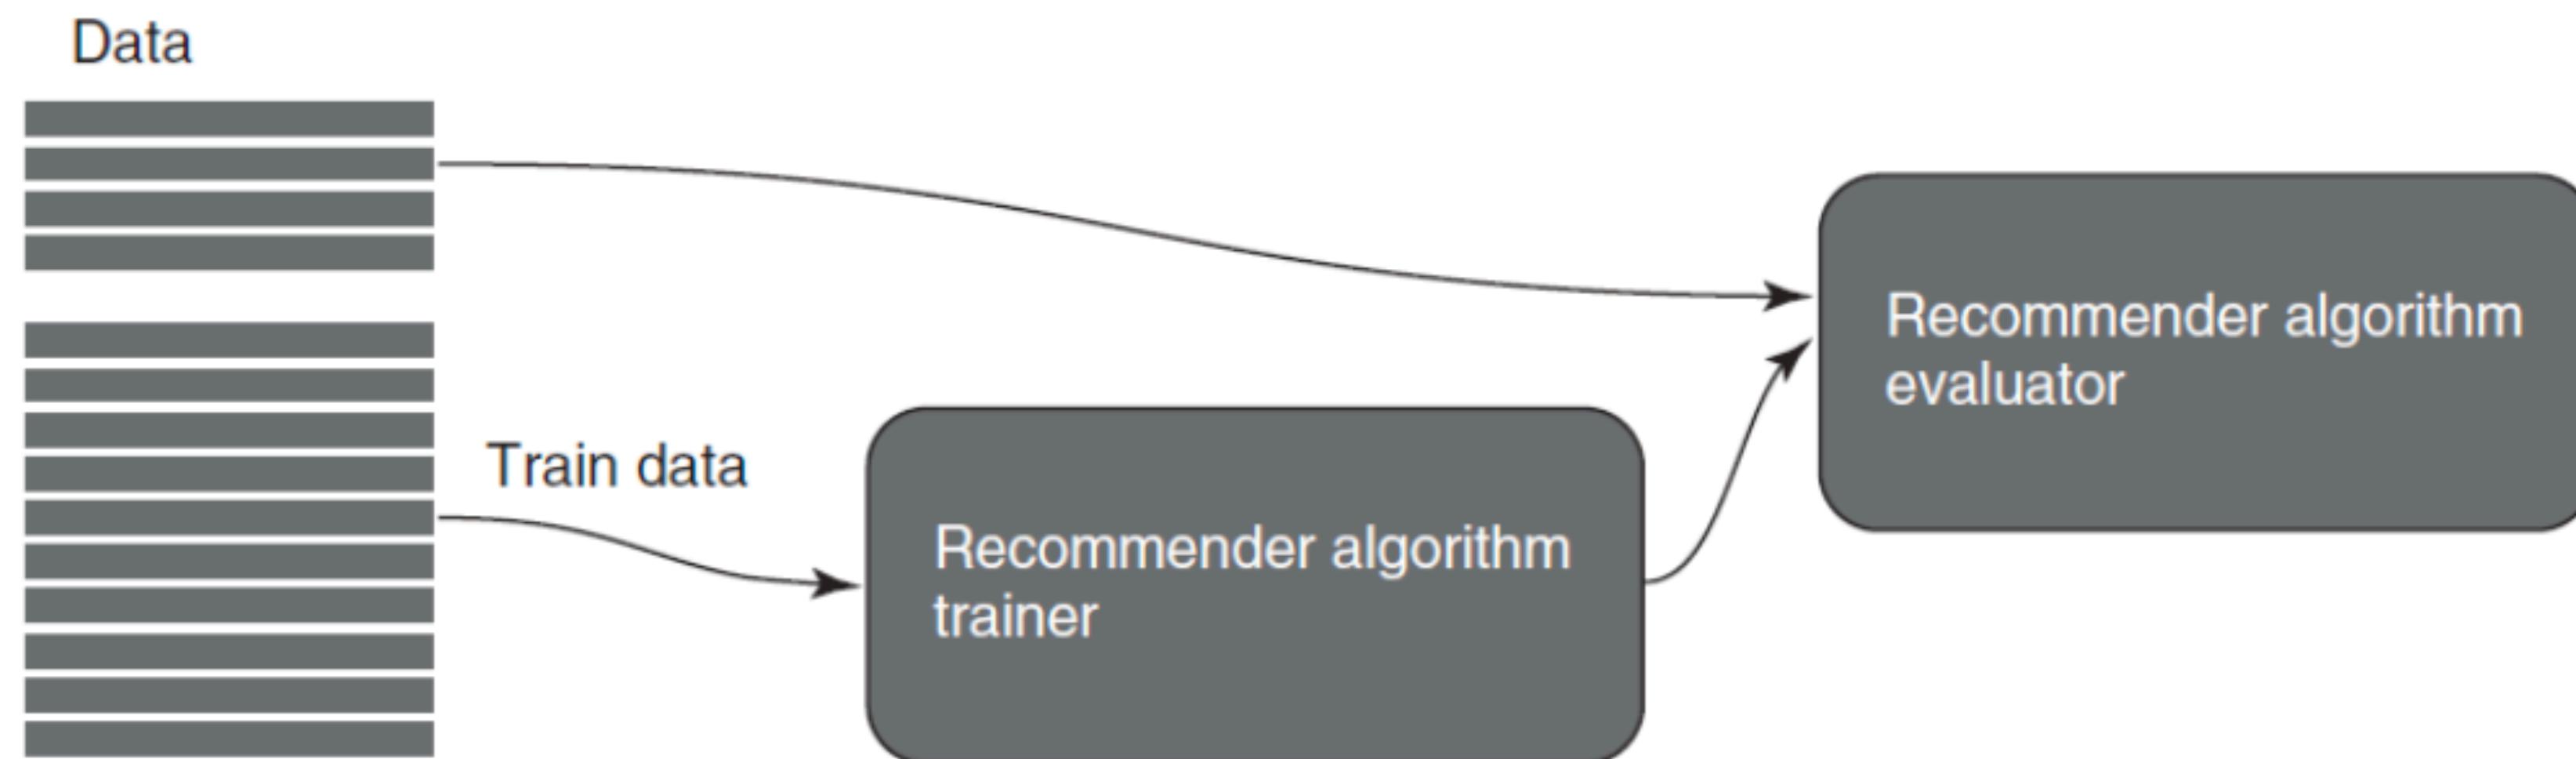
Evaluation

Rank Score

- DCG and NDCG: take into account the ‘position’ of the ranking list
- Hits at the beginning count more (more “gain”) and items of higher relevance are more important
- Given a rank position pos , and the graded relevance “rel” of an item i :
$$DCG = rel_1 + \sum_{i=1}^{pos} \frac{rel_i}{\ln(i)}$$
- Since some queries are harder than others and will likely produce lower DCG scores than easier queries, we normalize by scaling the results based on ideal results (i.e., $NDCG$)

Evaluation

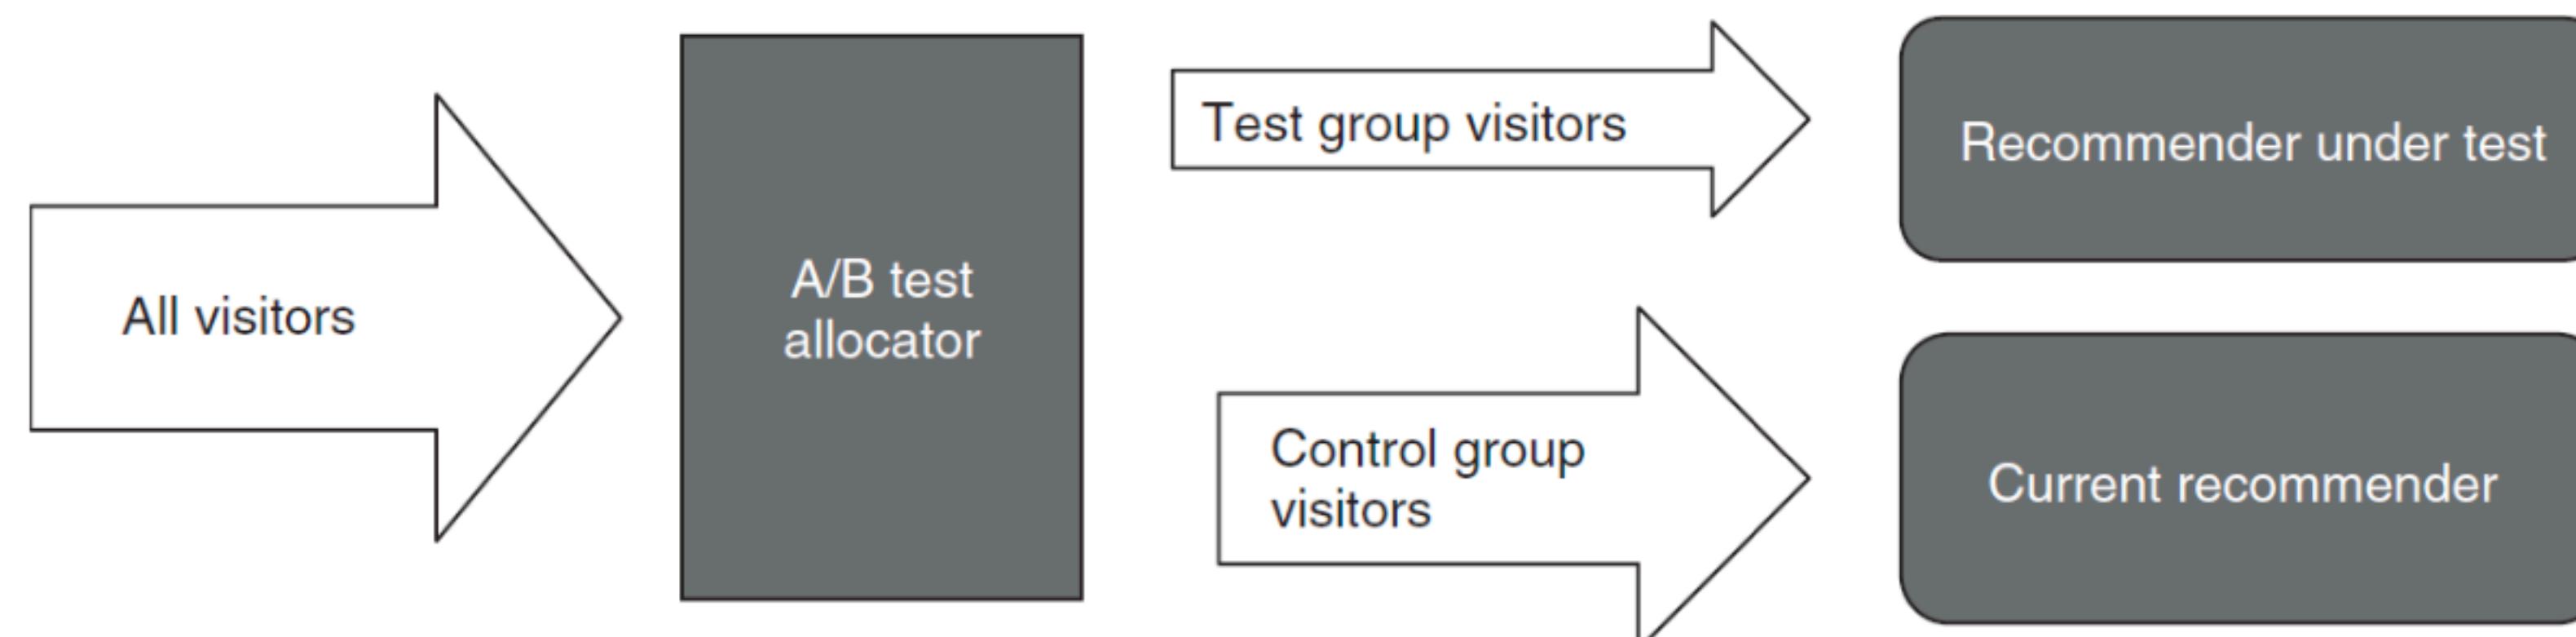
Offline evaluation



Evaluation

A/B Testing

- Offline evaluation shows good results
- Conduct A/B test to understand if there's a significant difference between the current recommender and the new one
- In practice, we divert a small percentage of traffic to the new feature
- Possible issue?



Summary

Summary

- Introduction to Recommender Systems
- Neighbourhood-based Collaborative Filtering
 - User-based
 - Item-based
- Evaluation

Thank You!